

The Cassowary Linear Arithmetic Constraint Solving Algorithm

GREG J. BADROS and ALAN BORNING

University of Washington

and

PETER J. STUCKEY

University of Melbourne

Linear equality and inequality constraints arise naturally in specifying many aspects of user interfaces, such as requiring that one window be to the left of another, requiring that a pane occupy the leftmost third of a window, or preferring that an object be contained within a rectangle if possible. Previous constraint solvers designed for user interface applications cannot handle simultaneous linear equations and inequalities efficiently. This is a major limitation, as such systems of constraints arise often in natural declarative specifications. We describe Cassowary—an incremental algorithm based on the dual simplex method, which can solve such systems of constraints efficiently. We have implemented the algorithm as part of a constraint-solving toolkit. We discuss the implementation of the toolkit, its application programming interface, and its performance.

Categories and Subject Descriptors: D.1.6 [**Software**]: Logic Programming; D.2.2 [**Software**]: Tools and Techniques; G.1.3 [**Mathematics of Computing**]: Numerical Linear Algebra; G.1.6 [**Mathematics of Computing**]: Optimization

General Terms: Algorithms, Design, Theory, Verification

Additional Key Words and Phrases: Cassowary, constraint-solving toolkit, constraints, user interface

1. INTRODUCTION

Linear equality and inequality constraints arise naturally in specifying many aspects of user interfaces, especially layout and other geometric relations. Inequality constraints, in particular, are needed to express relationships such as “inside,” “above,” “below,” “left-of,” “right-of,” and “overlaps.” For example, if we

This research was funded in part by a NSF Graduate Research Fellowship and a UW CSE Wilma Bradley Fellowship for Greg Badros, in part by NSF grant IIS-9975990, and in part by the Australian-American Educational Foundation (Fulbright Commission).

Authors' addresses: G. Badros and A. Borning, Dept. of Computer Science and Engineering, University of Washington, Box 352350, Seattle, WA 98195-2350; email: {badros,borning}@cs.washington.edu; P. Stuckey, Dept. of Computer Science and Software Engineering, University of Melbourne, Parkville, Victoria 3052, Australia; email: pjs@cs.mu.oz.au.

Permission to make digital/hard copy of all or part of this material without fee for personal or classroom use provided that the copies are not made or distributed for profit or commercial advantage, the ACM copyright/server notice, the title of the publication, and its date appear, and notice is given that copying is by permission of the ACM, Inc. To copy otherwise, to republish, to post on servers, or to redistribute to lists requires prior specific permission and/or a fee.

© 2001 ACM 1073-0516/01/1200-0267 \$5.00

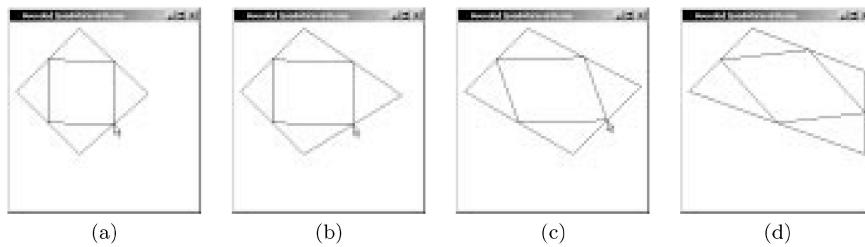


Fig. 1. Demonstrating a theorem about quadrilaterals.

are designing a Web document, we can express the requirement that `figure1` be to the left of `figure2` as the constraint `figure1.rightSide ≤ figure2.leftSide`.

It is important to be able to express preferences as well as requirements in a constraint system used for graphical layout. For example, we must be able to express a desire for stability when moving parts of an image: Things should stay where they were unless there is some reason for them to move. A second use for preferred constraints is to cope gracefully with invalid user inputs. For example, if the user tries to move a figure outside of its bounding window, it is reasonable for the figure just to bump up against the side of the window and stop, rather than causing an exception. A third use of nonrequired constraints is to balance conflicting desires, for example in laying out a graph.

In this article we describe Cassowary, an incremental constraint-satisfaction algorithm, that can solve such systems of constraints efficiently. A traditional demonstration of constraint-based graphics is the Quadrilateral theorem. Figure 1 shows four screen shots from our Smalltalk implementation of Cassowary. In the demonstration, each side of the quadrilateral is bisected, and lines are drawn between the midpoints. (These inner lines always form a parallelogram.) This is expressed as constraints that each midpoint lie halfway between the endpoints of its line. In addition, all points are constrained to be at least 10 pixels from the sides of the window—these additional inequality constraints make this collection of constraints beyond the scope of most interactive solvers.

In Figure 1(a) we picked up one of the midpoints with the mouse and began to move it by temporarily adding an *edit constraint* equating the position of the midpoint and the mouse. This constraint is strongly preferred but not required—we will violate it if necessary. The remaining points are weakly constrained to stay where they are by *stay constraints*. In Figure 1(b), the mouse has been moved to the right, and to keep the midpoint constraint satisfied the right-most vertex has moved as well. We continue moving to the right. In Figure 1(c), the rightmost vertex has run into the imaginary wall resulting from the constraint that it be at least 10 pixels from the window boundary, and can move no further, so the bottom vertex has begun moving. Finally, in Figure 1(d), the mouse has moved beyond the permitted region for the midpoint. The midpoint has moved as close to the mouse as possible, thus causing the two endpoints of its line to be pressed against the boundary as well.

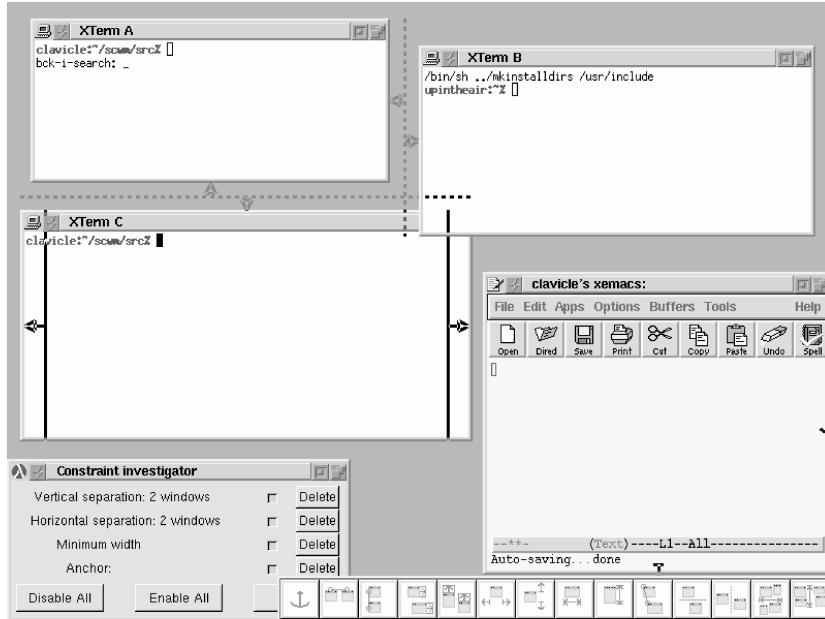


Fig. 2. Screen shot from the Scheme Constraints Window Manager.

Figure 2 illustrates a more complex application of Cassowary, namely SCWM, the Scheme Constraints Window Manager [Badros et al. 2000]. In the screen shot, XTerm A is constrained to be to the left of XTerm B, and above XTerm C. Additionally, XTerm C is required to have a minimum width, and the XEmacs window's southeast corner is anchored at its current location. All of these constraints are maintained as the user manipulates the windows. A button bar at the bottom of the screen provides a graphical interface for applying constraints to windows. In addition, we have turned on the “constraint investigator,” which provides a visual representation of the active constraints.

Efficient techniques are available for solving systems of linear constraints if the constraint network is acyclic. However, in trying to apply constraint solvers to real-world problems, we found that the collection of constraints was often cyclic, and included both equalities and inequalities. Cycles sometimes arise when the programmer added redundant constraints—the cycles *could* have been avoided by careful analysis. However, the analysis is an added burden on the programmer. Further, it is clearly contrary to the spirit of the whole enterprise to require programmers to be constantly on guard to avoid cycles and redundant constraints; after all, one of the goals in providing constraints is to allow programmers to state what relations they want to hold in a declarative fashion, leaving it to the underlying system to enforce these relations. For other applications, such as complex layout problems with conflicting goals, cycles seem unavoidable. Hence, a solver that can handle cycles of both equality and inequality constraints is highly desirable.

1.1 Constraint Hierarchies and Comparators

Since we want to express preferences as well as requirements in the constraint system, we need a specification for how conflicting preferences are to be traded off. *Constraint hierarchies* [Borning et al. 1992] provide a general theory for this. In a constraint hierarchy, each constraint has a strength. The required strength is special, in that required constraints must be satisfied. The other strengths all label nonrequired constraints. A constraint of a given strength completely dominates any constraint with a weaker strength—the strong constraint must be satisfied as well as possible before the weaker constraint can have any effect on the solution. In the theory, a *comparator* is used to compare different possible solutions to the constraints and select among them. Given a constraint hierarchy, ideally we would find one solution that is “better” than all others using this comparator. However, in most cases there will be many equally good solutions, and any one of these will be a correct answer.

A *labelled primitive constraint* is a primitive constraint labelled with a strength, written sc , where s is a strength and c is a primitive constraint. Strengths are non-negative integers; for clarity, we give them symbolic names. Strength 0, with the symbolic name required, is always reserved for required constraints. In the examples here, we use three strengths: required, strong, and weak. However, the approach permits any finite number of strengths.

Formally, a *constraint hierarchy* is a multiset of labelled primitive constraints. Given a constraint hierarchy H , H_0 denotes the multiset of required primitive constraints in H , with their labels removed. In the same way, we define the multisets H_1, H_2, \dots, H_n for strengths $1, 2, \dots, n$.

A *valuation* is a mapping of variables to values, written $\{x_1 \mapsto d_1, \dots, x_n \mapsto d_n\}$ where x_1, \dots, x_n are distinct variables and d_1, \dots, d_n are values (in this article, real numbers). The set S of *solutions* to the hierarchy is defined as follows. S_0 is the set of valuations such that all the H_0 constraints hold; from this we form S by eliminating all potential valuations that are worse than some other potential valuation using the comparator predicate *better*.

$$\begin{aligned} S_0 &= \{\theta \mid \forall c \in H_0 \ c\theta \text{ holds}\} \\ S &= \{\theta \mid \theta \in S_0 \wedge \forall \sigma \in S_0 \ \neg \text{better}(\sigma, \theta, H)\} \end{aligned}$$

Within this framework a number of variations are possible. One decision is whether we only compare solutions on a constraint-by-constraint basis (a *local* comparator), or whether we take some aggregate measure of the unsatisfied constraints of a given strength (a *global* comparator). A second choice is whether we are concerned only whether a constraint is satisfied or not (a *predicate* comparator), or whether we also want to know how nearly satisfied it is (a *metric* comparator). Constraints whose domain is a metric space, such as constraints on the reals, can have an associated metric error function. The error in satisfying such a constraint is 0 if and only if the constraint is satisfied, and becomes larger the less nearly satisfied the constraint is.

For inequality constraints it is important to use a metric rather than a predicate comparator [Borning et al. 1996]. Thus, plausible comparators for use with linear equality and inequality constraints are *locally-error-better*,

weighted-sum-better, and *least-squares-better*. For a given collection of constraints, Cassowary finds a weighted-sum-better solution. (Every weighted-sum-better solution is also a locally-error-better solution).

We assume an error function $e(c\theta)$ that returns a nonnegative real number indicating how nearly the primitive constraint c is satisfied for a valuation θ . For a linear equation c of the form $l = 0$ (l is a linear expression), the error function is $e(c\theta) = |l\theta|$, while for a linear inequality c of the form $l \leq 0$, the error function is

$$e(c\theta) = \begin{cases} l\theta & \text{where } l\theta > 0 \\ 0 & \text{otherwise.} \end{cases}$$

For the weighted-sum-better comparator, a solution θ can be thought of as mapping a constraint hierarchy H to a tuple of errors defined by

$$[\Sigma_{c \in H_1} e(c\theta), \Sigma_{c \in H_2} e(c\theta), \dots, \Sigma_{c \in H_n} e(c\theta)].$$

Then one solution θ is better than another solution σ if the corresponding tuple of errors is lexicographically smaller.

For example, consider the constraint hierarchy

$$\begin{aligned} \text{required } 2x_m &= x_l + x_r \\ \text{strong } x_r &= 90 \\ \text{weak } x_l &= 50 \\ \text{weak } x_r &= x_m + 10 \end{aligned}$$

and consider the two solutions $\theta_1 \equiv \{x_l \mapsto 50, x_m \mapsto 60, x_r \mapsto 70\}$ and $\theta_2 \equiv \{x_l \mapsto 50, x_m \mapsto 70, x_r \mapsto 90\}$. The corresponding error tuples are $[20, 0]$ and $[0, 10]$, so θ_2 is weighted-sum-better than θ_1 , since $[0, 10]$ is lexicographically less than $[20, 0]$.

In what follows we consider only nonrequired constraints of the form strength $v = d$ where v is a variable and d a value. Note that we can straightforwardly translate any constraint hierarchy into an equivalent constraint hierarchy in this form by replacing a nonrequired constraint strength $l = 0$ by required $l = e_l$, strength $e_l = 0$, and replacing a nonrequired constraint strength $l \geq 0$ by required $l - e_l \geq 0$, required $e_l \geq 0$, strength $e_l = 0$.

1.2 Adapting the Simplex Algorithm

Linear programming is concerned with solving the following problem:

Consider a collection of n real-valued variables x_1, \dots, x_n , each of which is constrained to be nonnegative: $x_i \geq 0$ for $1 \leq i \leq n$. Suppose there are m linear equality or inequality constraints over the x_i , each of the form:

$$\begin{aligned} a_1x_1 + \dots + a_nx_n &= c, \\ a_1x_1 + \dots + a_nx_n &\leq c, \text{ or} \\ a_1x_1 + \dots + a_nx_n &\geq c. \end{aligned}$$

Given these constraints, find values for the x_i that minimize (or maximize) the value of the *objective function* $e + d_1x_1 + \dots + d_nx_n$.

This problem has been heavily studied for the past fifty years. The most commonly used technique for solving it is the simplex algorithm, developed by Dantzig in the 1940s. It now has numerous variations, but, unfortunately, existing implementations of the simplex algorithm are not readily usable for user interface applications.

The principal difficulty is incrementality. For interactive graphical applications, we need to solve similar problems repeatedly, rather than solving a single problem once. To achieve interactive response times, fast incremental algorithms that exploit prior computations are needed. There are two common cases that algorithmic changes should try to improve. First, when moving an object with a mouse or other input device, we typically represent this interaction as a one-way constraint relating the mouse position to the desired x and y coordinates of a part of the figure. For each screen refresh, we must resatisfy the same collection of constraints while varying only the mouse location input. The second common operation that incremental algorithms can optimize occurs when editing an object in a complex system. Ideally, when adding or removing a small number of constraints, we would like to avoid resolving the entire system. Although the performance requirements for this case are less stringent than for the first case, since it is a less frequent operation, we still wish to increase performance by reusing as much of the previous solution as possible.

Another important issue when applying simplex to user interface applications is defining a suitable objective function. We must accommodate nonrequired constraints of different strengths, which is analogous to multiobjective linear programming problems. Also, the objective function in the standard simplex algorithm must be a linear expression; but the objective functions for the locally-error-better, weighted-sum-better, and least-squares-better comparators are all nonlinear. For Cassowary, we avoid the least-squares-better comparator and use a quasi-linear objective function to implement the weighted-sum-better comparator (see Section 2.3).

Finally, a minor issue is accommodating variables that may take on both positive and negative values, which is generally the case in user interface applications. (The standard simplex algorithm requires all variables to be nonnegative). Here we adopt efficient techniques developed for implementing constraint logic programming languages (see Section 2.1).

1.3 Implementation Overview

The Cassowary algorithm has been implemented in Smalltalk, C++, and Java as part of our Cassowary Constraint Solving Toolkit, and is freely available for download [Badros and Borning 2001]. The library performs surprisingly well, and a summary of our performance measurements is given in Section 3. A reimplemention of the algorithm based on this article is also reasonable, given a knowledge of the simplex algorithm. The electronic appendix to this article, available via the ACM Digital Library, provides additional details regarding our implementations of the Cassowary algorithm.

1.4 Related Work

There is a long history of using constraints in user interfaces and interactive systems, beginning with Ivan Sutherland's pioneering Sketchpad system [Sutherland 1963]. Most of the current systems use one-way constraints (e.g., Hudson and Smith [1996]; Myers [1996]), or local propagation algorithms for acyclic collections of multiway constraints (e.g., Sannella et al. [1993]; Vander Zanden [1996]). Indigo [Borning et al. 1996] handles acyclic collections of inequality constraints, but not cycles (simultaneous equations and inequalities). User interface systems that handle simultaneous linear equations include DETAiL [Hosobe et al. 1996] and Ultraviolet [Borning and Freeman-Benson 1998]. A number of researchers (including the second author) have experimented with a straightforward use of a simplex package in a UI constraint solver, but the speed was not satisfactory for interactive use.

Baraff [1994] describes a quadratic optimization algorithm for solving linear constraints that arise in modeling physical systems. Finally, much of the work on constraint solvers has been in the logic programming and constraint logic programming communities. Current constraint logic programming languages such as CLP(\mathcal{R}) [Jaffar et al. 1992] include efficient solvers for linear equalities and inequalities (see Marriott and Stuckey [1998] for a survey). However, these solvers use a refinement model of computation, in which the values determined for variables are successively refined as the computation progresses, but there is no specific notion of state and change. As a result, these systems are not especially well suited for building interactive graphical applications.

An earlier article [Borning et al. 1997] describes both the original version of Cassowary (in much less detail than in this article), and also the related QOCA algorithm. (QOCA uses much the same solving technique as Cassowary, but implements a least-squares-better comparator, which strongly penalizes outlying values when comparing constraints of the same strength.) The incremental constraint deletion procedure is described in Huynh and Marriot [1995].

2. INCREMENTAL SIMPLEX

As you see, the subject of linear programming is surrounded by nomenclatural and terminological thickets. Both of these thorny defenses are lovingly cultivated by a coterie of stern acolytes who have devoted themselves to the field. Actually, the basic ideas of linear programming are quite simple. —*Numerical Recipes* [Press et al. 1989, page 424].

We now describe an incremental version of the simplex algorithm, adapted for our Cassowary algorithm for interactive graphical applications. In the description we use a running example, illustrated by the diagram in Figure 3.

The constraints on the variables in Figure 3 are as follows: x_m is constrained to be the midpoint of the line from x_l to x_r , and x_l is constrained to be at least 10 to the left of x_r . All variables must lie in the range -10 to 100 . (To keep the presentation manageable, we deal only with the x coordinates. Adding analogous constraints on the y coordinates is straightforward, but would double the number of constraints in our example.) Since $x_l < x_m < x_r$ in any solution, we

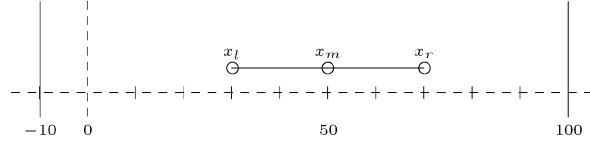


Fig. 3. Simple constrained picture.

simplify the problem by removing the redundant bounds constraints. However, even with these simplifications, the resulting constraints have a cyclic constraint graph and cannot be handled by methods such as Indigo [Bornig et al. 1996].

With the simplifications, the constraints are

$$\begin{aligned} 2x_m &= x_l + x_r \\ x_l + 10 &\leq x_r \\ x_l &\geq -10 \\ x_r &\leq 100 \end{aligned}$$

2.1 Augmented Simplex Form

Suppose we wish to minimize the distance between x_m and x_l , or in other words, minimize $x_m - x_l$. (This simple objective function is used as an initial example only; an objective function for an interactive manipulation problem is described in Section 2.3).

The basic simplex algorithm does not itself handle variables that may take negative values (so-called *unrestricted variables*). It instead imposes an implicit constraint $x \geq 0$ on all variables occurring in its equations. Augmented simplex form allows us to handle unrestricted variables efficiently and simply; it was developed for implementing constraint logic programming languages [Marriott and Stuckey 1998], and we have adopted it here. Conceptually it uses *two* tableaux rather than one. All of the unrestricted variables from the original constraints C will be placed in C_U , the unrestricted variable tableau. C_S , the simplex tableau, contains only variables constrained to be nonnegative (the *restricted variables*).

Thus, an optimization problem is in *augmented simplex form* if the constraints C have the form $C_U \wedge C_S \wedge C_I$ where C_U and C_S are conjunctions of linear arithmetic equations, C_I is $\bigwedge \{x \geq 0 \mid x \in vars(C_S)\}$, and the objective function f is a linear expression over variables in C_S .

The simplex algorithm is used to determine an optimal solution for the equations in the simplex tableau C_S , ignoring the unrestricted variable tableau C_U during the optimization procedure. (C_U need only be considered when finding the feasible region, prior to optimization.) The equations in the C_U are then used to determine values for the unrestricted variables.

It is not difficult to rewrite an arbitrary optimization problem over linear real equations and inequalities into augmented simplex form. The first step is to convert inequalities to equations. Each inequality of the form $l \geq 0$, where l is a linear real expression, can be replaced with $l = s \wedge s \geq 0$ where s is a newly-introduced nonnegative *slack* variable.

For example, the constraints for Figure 3 can be rewritten as

minimize $x_m - x_l$ subject to

$$\begin{array}{l} 2x_m = x_l + x_r \\ x_l + 10 + s_1 = x_r \\ x_l - s_2 = -10 \\ x_r + s_3 = 100 \\ 0 \leq s_1, s_2, s_3 \end{array}$$

We now separate the equalities into C_U and C_S . Initially all equations are in C_S . We move the unrestricted variables into C_U using Gauss-Jordan elimination. To do this, we select an equation in C_S containing an unrestricted variable x_u and remove the equation from C_S . We then solve the equation for x_u , yielding a new equation $x_u = l$ for some linear expression l . We then substitute l for all remaining occurrences of x_u in C_S , C_U , and f , and place the equation $x_u = l$ in C_U . The process is repeated until there are no more unrestricted variables in C_S . In our example, $x_r + s_3 = 100$ can be used to substitute $100 - s_3$ for x_r , yielding

minimize $x_m - x_l$ subject to

$$\begin{array}{rcl} x_r = 100 - s_3 & C_U \\ \hline 2x_m = x_l + 100 - s_3 & C_S \\ x_l + 10 + s_1 = 100 - s_3 \\ x_l - s_2 = -10 \\ 0 \leq s_1, s_2, s_3 & C_I \end{array}$$

Next, the first equation of C_S can be used to substitute $50 + \frac{1}{2}x_l - \frac{1}{2}s_3$ for x_m , giving, minimize $50 - \frac{1}{2}x_l - \frac{1}{2}s_3$ subject to

$$\begin{array}{rcl} x_m = 50 + \frac{1}{2}x_l - \frac{1}{2}s_3 \\ x_r = 100 - s_3 \\ \hline x_l + 10 + s_1 = 100 - s_3 & C_U \\ x_l - s_2 = -10 \\ 0 \leq s_1, s_2, s_3 & C_I \end{array}$$

Now we move x_l to C_U using $x_l = s_2 - 10$, giving, minimize $55 - \frac{1}{2}s_2 - \frac{1}{2}s_3$ subject to

$$\begin{array}{rcl} x_m = 45 + \frac{1}{2}s_2 - \frac{1}{2}s_3 \\ x_r = 100 - s_3 \\ \hline x_l = s_2 - 10 & C_U \\ s_2 + s_1 = 100 - s_3 & C_S \\ 0 \leq s_1, s_2, s_3 & C_I \end{array}$$

The general algorithm is shown in Figure 4. Note that in this and subsequent algorithms we treat constraints modulo logical equivalence, so that, for example, the constraint $2x_m - x_l = x_r$ is identical to the constraint $x_m = \frac{1}{2}x_r + \frac{1}{2}x_l$, and use this implicitly in matching the forms of constraints.

```

augmented-simplex-form( $C, f$ )
 $C_I := \emptyset; C_U := \emptyset;$ 
for each inequality  $l \geq 0$  in  $C$ 
     $C := C - \{l \geq 0\}$ 
    let  $s$  be a new slack variable
     $C := C \cup \{l = s\}$ 
     $C_I := C_I \cup \{s \geq 0\}$ 
endfor
while exists variable  $y$  in  $\text{vars}(C) - \text{vars}(C_I)$ 
    let  $y = l \in C$  where  $y \notin \text{vars}(l)$ 
     $C := C - \{c\}$ 
    replace  $y$  by  $l$  in  $C_U \cup C$  and in  $f$ 
     $C_U := C_U \cup \{y = l\}$ 
endwhile
 $C_S := C$ 
return  $(C_U, C_S, C_I, f)$ 

```

Fig. 4. Take an initial constraint C and objective function f and return an equivalent augmented simplex form $C_U \cup C_S \cup C_I$ with f .

Hereafter, the labels for C_U and C_S are omitted: constraints above the horizontal line are in C_U , and constraints below the line are in C_S . Also, C_I will be omitted entirely—any variable occurring below the horizontal line is implicitly constrained to be nonnegative.

The simplex method works by taking an optimization problem in “basic feasible solved form” (a type of normal form) and repeatedly applying matrix operations to obtain new basic feasible solved forms. Once we have split the equations into C_U and C_S , we can ignore C_U for purposes of optimization.

In the implementation, all variables that may be accessed from outside the solver are unrestricted. Only error or slack variables are represented as restricted variables, and these variables occur only within the solver (see the electronic appendix for further details). The primary benefit of this simplification is that the programmer using the solver always uses just the one kind of variable. A minor benefit is that only the external, unrestricted variables actually store their values as a field in the variable object; the values of restricted variables are just given by the tableau. A minor drawback is that the constraint $x \geq 0$ must be represented explicitly for an external variable x . (In any case, $x \geq b$ would need to be represented explicitly for any other constant $b \neq 0$).

In our running example, the constraints imply that x_r is nonnegative. However, since x_r is accessible from outside the solver, we represent it as unrestricted. This does not change the solutions found. Also, we show the operations as modifying C_U as well as C_S . It would be possible to modify just C_S and leave C_U unchanged, using C_U only to define values for the variables on the left-hand side of its equations. This would speed up pivoting, but it would make the incremental updates of the constants in edit constraints slower (Section 2.4). Because the latter is a much more frequent operation, we do actually modify both C_U and C_S in the implementation.

An augmented simplex form optimization problem is in *basic-feasible-solved form* if the equations are of the form

$$x_0 = c + a_1x_1 + \cdots + a_nx_n,$$

where the variable x_0 does not occur in any other equation or in the objective function. If the equation is in C_S , c must be nonnegative. However, there is no such restriction on the constants for the equations in C_U . In either case, the variable x_0 is said to be *basic* and the other variables in the equation are *parameters*. A problem in basic-feasible-solved form defines a *basic-feasible solution*, which is obtained by setting each parametric variable to 0 and each basic variable to the value of the constant in the right-hand side. Note that the *basic-feasible solution* may not be an optimal solution.

For instance, the following constraint is in basic-feasible-solved form, and is equivalent to the problem above.

minimize $55 - \frac{1}{2}s_2 - \frac{1}{2}s_3$ subject to

$$\begin{aligned}x_l &= -10 + s_2 \\x_m &= 45 + \frac{1}{2}s_2 - \frac{1}{2}s_3 \\x_r &= 100 - s_3 \\ \hline s_1 &= 100 - s_2 - s_3\end{aligned}$$

The basic feasible solution corresponding to this basic feasible solved form is

$$\{x_l \mapsto -10, x_m \mapsto 45, x_r \mapsto 100, s_1 \mapsto 100, s_2 \mapsto 0, s_3 \mapsto 0\}.$$

The value of the objective function with this solution is 55.

2.2 Simplex Optimization

We now describe how to find an optimum solution to a constraint in basic-feasible-solved form. Except for the operations on the additional unrestricted variable tableau C_U , the material presented in this section is simply phase II of the standard two-phase simplex algorithm.

The simplex algorithm finds the optimum by repeatedly looking for an “adjacent” basic-feasible-solved form whose basic feasible solution decreases the value of the objective function that we are minimizing. When no such adjacent basic feasible solved form can be found, we have achieved an optimum. The underlying operation is called *pivoting*, and involves exchanging a basic and a parametric variable using matrix operations. Thus, by “adjacent” we mean the new basic-feasible-solved form can be reached by performing a single pivot.

In our example, increasing s_2 from 0 will decrease the value of the objective function we are minimizing. We must be careful: We cannot increase the value of s_2 indefinitely, as this may cause the value of some other basic nonnegative variable to become negative. We must examine the equations in C_S . The equation $s_1 = 100 - s_2 - s_3$ allows s_2 to take at most a value of 100, because if s_2 becomes larger than this, then s_1 would become negative. The equations above the horizontal line do not restrict s_2 , since whatever value s_2 takes, the unrestricted variables x_l and x_m can take values to satisfy the equations. In general, we choose the most restrictive equation in C_S and use it to eliminate s_2 . In case of ties, we arbitrarily break the tie. In this example, the most restrictive equation (there is only one) is $s_1 = 100 - s_2 - s_3$. Writing s_2 as the subject, we obtain $s_2 = 100 - s_1 - s_3$. We replace s_2 everywhere by $100 - s_1 - s_3$, and obtain

minimize $5 + \frac{1}{2}s_1$ subject to

$$\begin{aligned} x_l &= 90 - s_1 - s_3 \\ x_m &= 95 - \frac{1}{2}s_1 - s_3 \\ x_r &= 100 - s_3 \\ \hline s_2 &= 100 - s_1 - s_3 \end{aligned}$$

We have just performed a pivot, having moved s_1 out of the set of basic variables and replaced it by moving s_2 into the basis.

We continue this process. Increasing the value of s_1 would increase the value of the objective function (which we are trying to minimize, so we don't want to do this). Note that decreasing s_1 would decrease the objective function's value, but as s_1 is constrained to be nonnegative, it already takes its minimum value of 0 in the associated basic feasible solution. Hence, we are at an optimal solution.¹

In general, the simplex algorithm applied to C_S is described as follows. We are given a problem in basic-feasible-solved form in which the variables x_1, \dots, x_n are basic and the variables y_1, \dots, y_m are parameters.

minimize $e + \sum_{j=1}^m d_j y_j$ subject to

$$\begin{aligned} \bigwedge_{i=1}^n x_i &= c_i + \sum_{j=1}^m a_{ij} y_j \quad \wedge \\ \bigwedge_{i=1}^n x_i &\geq 0 \quad \wedge \quad \bigwedge_{j=1}^m y_j \geq 0. \end{aligned}$$

Select an entry variable y_J such that $d_J < 0$. (An entry variable is one that will enter the basis, i.e., it is currently parametric and we want to make it basic). Pivoting on such a variable can only decrease the value of the objective function. If no such variable exists, the optimum has been reached. Now determine the exit variable x_I . We must choose this variable so that it maintains basic-feasible-solved form by ensuring that the new c_i 's are still positive after pivoting. That is, we must choose an x_I so that $-c_I/a_{IJ}$ is a minimum element of the set

$$\{-c_i/a_{ij} \mid a_{ij} < 0 \text{ and } 1 \leq i \leq n\}.$$

If there were no i for which $a_{ij} < 0$, then we could stop, since the optimization problem would be unbounded and so would not have a minimum: We could choose y_J to take an arbitrarily large value, and thus make the objective function arbitrarily small. However, this is not an issue in our context, since our optimization problems will always have a lower bound of 0.

We proceed to choose x_I , and pivot x_I out and replace it with y_J to obtain the new-basic-feasible solution. We continue this process until an optimum is reached. The algorithm is specified in Figure 5, and takes as inputs the simplex tableau C_S and the objective function f .

Finally, we must consider the problem of cycling. A basic feasible solution is *degenerate* if some of the basic variables take the value 0. If this is the case, we could perform a pivot that does not decrease the value of the objective function and that does not change the corresponding basic feasible solution. In such a

¹If we had an unrestricted variable in the objective function, the optimization would be unbounded. This possibility is not an issue for our algorithm due to the nature of the objective functions that arise from hierarchical constraints (see Section 2.3).

```

simplex_opt( $C_S, f$ )
repeat
    % Choose variable  $y_J$  to become basic
    if for each  $j \in \{1, \dots, m\}$   $d_j \geq 0$  then
        return ( $C_S, f$ ) % an optimal solution has been found
    endif
    choose  $J \in \{1, \dots, m\}$  such that  $d_J < 0$ 
    % Choose variable  $x_I$  to become non-basic
    choose  $I \in \{1, \dots, n\}$  such that
         $-c_I/a_{IJ} = \min_{i \in \{1, \dots, n\}} \{-c_i/a_{ij} \mid a_{ij} < 0\}$ 
     $l := (x_I - c_I - \sum_{j=1, j \neq J}^m a_{Ij}y_j)/a_{IJ}$ 
    replace  $y_J$  by  $l$  in  $f$ 
    for each  $i \in \{1, \dots, n\}$ 
        if  $i \neq I$  then replace  $y_J$  by  $l$  in row  $i$  endif
    endfor
    replace the  $I^{th}$  row by  $y_J = l$ 
endrepeat

```

Fig. 5. Simplex optimization of a system of equations C_S with objective function f in basic-feasible-solution form.

case, there is a danger of pivoting back to the original problem. In practice, such cycling does not occur very often, and in any event there are simple methods to avoid it. One of these is Bland's anti-cycling rule, which we use in our implementations. We number the variables (either with their hash code or a unique ID). When selecting a variable to enter the basis, we choose the lowest numbered variable. (This is done in the line “choose $J \in \{1, \dots, m\}$ such that $d_J < 0$ ” in Figure 5.) Similarly, when selecting a variable to exit the basis, if a tie occurs (next line), we choose the lowest numbered basic variable to break the tie.

2.3 Handling Nonrequired Constraints

Suppose the user wishes to edit x_m in the diagram and have x_l and x_r weakly stay where they are. This adds the nonrequired constraints *edit* x_m , *stay* x_l , and *stay* x_r . Suppose further that we are trying to move x_m to position 50, and that x_l and x_r are currently at 30 and 60, respectively. We are thus imposing the constraints strong $x_m = 50$, weak $x_l = 30$, and weak $x_r = 60$.

To form an objective function for the weighted-sum-better comparator, we sum the errors for each constraint, weighting the errors so that satisfying any strong constraint is always more important than satisfying any combination of weaker constraints. To implement this weighting, we might try using large coefficients, e.g., making the errors for strong constraints 1000 times larger than those for weak constraints. However, for sufficiently large values for the constrained variables, we might nevertheless end up in a situation in which weak constraints were satisfied in preference to strong ones.

Instead, to ensure that strong constraints are *always* satisfied in preference to weak ones, Cassowary uses symbolic weights for the coefficients in the objective function, rather than real numbers. These symbolic weights are represented as tuples and ordered lexicographically. In the presentation that follows,

we write these symbolic weights as pairs, such as [1, 2], which represents the symbolic weight consisting of the unit weight for the strong strength plus twice the unit weight for the weak strength.

For our example, the objective function is

$$[1, 0] \times |x_m - 50| + [0, 1] \times |x_l - 30| + [0, 1] \times |x_r - 60|$$

where [1, 0] is the weight (in lexicographic numbering) of a strong constraint and [0, 1] is the weight of a weak constraint. Due to the absolute value operators, this objective function is not linear, and hence the simplex method is not applicable directly. We now show how we can solve the problem using *quasi-linear optimization*.

Both the edit and the stay constraints are represented as equations of the form

$$v = \alpha + \delta_v^+ - \delta_v^-$$

where δ_v^+ and δ_v^- are nonnegative variables representing the deviation of v from the desired value α . If the constraint is satisfied, both δ_v^+ and δ_v^- will be 0. Otherwise δ_v^+ will be positive and δ_v^- will be 0 if v is too big, or vice versa if v is too small.² Because we want δ_v^+ and δ_v^- to be 0 if possible, we make them part of the objective function, with larger coefficients for the error variables of stronger constraints. (We need to use the pair of variables to satisfy simplex's nonnegativity restriction, since the variables δ_v^+ and δ_v^- will be part of the objective function).

Translating the constraints strong $x_m = 50$, weak $x_l = 30$, and weak $x_r = 60$, which arise from the edit and stay constraints, we obtain

$$\begin{aligned} x_m &= 50 + \delta_{x_m}^+ - \delta_{x_m}^- \\ x_l &= 30 + \delta_{x_l}^+ - \delta_{x_l}^- \\ x_r &= 60 + \delta_{x_r}^+ - \delta_{x_r}^- \\ 0 &\leq \delta_{x_m}^+, \delta_{x_m}^-, \delta_{x_l}^+, \delta_{x_l}^-, \delta_{x_r}^+, \delta_{x_r}^- \end{aligned}$$

as well as the original constraints:

$$\begin{aligned} 2x_m &= x_l + x_r \\ x_l + 10 &\leq x_r \\ x_l &\geq -10 \\ x_r &\leq 100 \end{aligned}$$

The objective function for our example can now be restated as

$$\text{minimize } [1, 0]\delta_{x_m}^+ + [1, 0]\delta_{x_m}^- + [0, 1]\delta_{x_l}^+ + [0, 1]\delta_{x_l}^- + [0, 1]\delta_{x_r}^+ + [0, 1]\delta_{x_r}^-.$$

An optimal solution of this problem can now be found using the simplex algorithm. The only point to notice in using symbolic weights as coefficients is that the comparison $d_J < 0$ choosing an entry variable J is a lexicographic comparison. The optimization results in a tableau:

²Although the equation may be satisfied with both δ_v^+ and δ_v^- nonzero, the simplex optimization itself forces at least one of them to be zero (see Section 2.4).

minimize $[0, 10] + [1, 2]\delta_{x_m}^+ + [1, -2]\delta_{x_m}^- + [0, 2]\delta_{x_l}^- + [0, 2]\delta_{x_r}^-$ subject to

$$\begin{array}{rccccc}
 x_m & = & 50 & +\delta_{x_m}^+ & -\delta_{x_m}^- & \\
 x_l & = & 30 & & +\delta_{x_l}^+ & -\delta_{x_l}^- \\
 x_r & = & 70 & +2\delta_{x_m}^+ & -2\delta_{x_m}^- & -\delta_{x_l}^+ & +\delta_{x_l}^- \\
 \hline
 s_1 & = & 30 & +2\delta_{x_m}^+ & -2\delta_{x_m}^- & -2\delta_{x_l}^+ & +2\delta_{x_l}^- \\
 s_3 & = & 30 & -2\delta_{x_m}^+ & +2\delta_{x_m}^- & +\delta_{x_l}^+ & -\delta_{x_l}^- \\
 \delta_{x_r}^+ & = & 10 & +2\delta_{x_m}^+ & -2\delta_{x_m}^- & -\delta_{x_l}^+ & +\delta_{x_l}^- +\delta_{x_r}^- \\
 s_2 & = & 40 & & +\delta_{x_l}^+ & -\delta_{x_l}^- &
 \end{array}$$

This corresponds to the solution $\{x_m \mapsto 50, x_l \mapsto 30, x_r \mapsto 70\}$ illustrated in Figure 3. Notice that the weak stay constraint on x_r is not satisfied ($\delta_{x_r}^+$ is nonzero, read directly from the second to last line of the above tableau).

In general, the objective function corresponding to a constraint hierarchy H is

$$[1, 0, \dots, 0]\Sigma_{\delta_1 \in \Delta_1}\delta_1 + [0, 1, \dots, 0]\Sigma_{\delta_2 \in \Delta_2}\delta_2 + \dots + [0, 0, \dots, 1]\Sigma_{\delta_n \in \Delta_n}\delta_n,$$

where Δ_i is the set of error variables for constraints of strength i . Note that the minimum of this function is clearly bounded below, since all error variables are nonnegative.

2.4 Incrementality: Resolving the Optimization Problem

Now suppose the user moves the mouse (which is editing x_m) to $x = 60$. We wish to solve a new problem, with constraints strong $x_m = 60$, weak $x_l = 30$, and weak $x_r = 70$ (so that x_l and x_r should stay where they are if possible).

There are two steps. First, we modify the tableau to reflect the new constraints we wish to solve. Second, we resolve the optimization problem for this modified tableau.

Let us first examine how to modify the tableau to reflect the new values of the stay constraints. This will not require reoptimizing the tableau, since we know that the new stay constraints are satisfied exactly. Suppose the previous stay value for variable v was α , and in the current solution v takes value β . The current tableau contains the information that

$$v = \alpha + \delta_v^+ - \delta_v^-,$$

and we need to modify this so that instead,

$$v = \beta + \delta_v^+ - \delta_v^-.$$

There are two cases to consider: (a) both δ_v^+ and δ_v^- are parameters, or (b) one of them is basic.

In case (a), v must take the value α in the current solution since both δ_v^+ and δ_v^- take the value 0, and

$$v = \alpha + \delta_v^+ - \delta_v^-.$$

Hence, $\beta = \alpha$ and no changes need to be made.

In case (b), assume without loss of generality that δ_v^+ is basic. Now, in the original equation representing the stay constraint, the coefficient for δ_v^+ is the negative of the coefficient for δ_v^- . Every constraint in the simplex tableau is simply a linear sum of the original constraints, since the only modifications to the tableau are by pivoting. Since the δ_v^+ and δ_v^- variables occur in no other constraints, this relation between the coefficients will continue to hold as we perform pivots. In other words, δ_v^+ and δ_v^- come in pairs: any equation that contains δ_v^+ will also contain δ_v^- (with the coefficient negated), and vice versa. Since δ_v^+ is assumed to be basic, it occurs exactly once in an equation with constant c , and further this equation also contains the only occurrence of δ_v^- . In the current solution,

$$\{v \mapsto \beta, \delta_v^+ \mapsto c, \delta_v^- \mapsto 0\}$$

and, since the equation

$$v = \alpha + \delta_v^+ - \delta_v^-$$

holds, $\beta = \alpha + c$. To replace the equation

$$v = \alpha + \delta_v^+ - \delta_v^-$$

by

$$v = \beta + \delta_v^+ - \delta_v^-$$

we simply need to replace the constant c in the row for δ_v^+ by 0. Since there are no other occurrences of δ_v^+ and δ_v^- we have replaced the old equation with the new one.

For our example, to update the tableau for the new values for the stay constraints on x_l and x_r we simply set the constant of the second to last equation (the equation for $\delta_{x_r}^+$) to 0. The tableau now is

minimize $[0, 0] + [1, 2]\delta_{x_m}^+ + [1, -2]\delta_{x_m}^- + [0, 2]\delta_{x_l}^- + [0, 2]\delta_{x_r}^-$ subject to

$x_m = 50$	$+ \delta_{x_m}^+$	$- \delta_{x_m}^-$			
$x_l = 30$			$+ \delta_{x_l}^+$	$- \delta_{x_l}^-$	
$x_r = 70$	$+ 2\delta_{x_m}^+$	$- 2\delta_{x_m}^-$	$- \delta_{x_l}^+$	$+ \delta_{x_l}^-$	
$s_1 = 30$	$+ 2\delta_{x_m}^+$	$- 2\delta_{x_m}^-$	$- 2\delta_{x_l}^+$	$+ 2\delta_{x_l}^-$	
$s_3 = 30$	$- 2\delta_{x_m}^+$	$+ 2\delta_{x_m}^-$	$+ \delta_{x_l}^+$	$- \delta_{x_l}^-$	
$\delta_{x_r}^+ = 0$	$+ 2\delta_{x_m}^+$	$- 2\delta_{x_m}^-$	$- \delta_{x_l}^+$	$+ \delta_{x_l}^-$	$+ \delta_{x_r}^-$
$s_2 = 40$			$+ \delta_{x_l}^+$	$- \delta_{x_l}^-$	

(For completeness, we update the constant part of the objective function, as well as its other terms, in our running example. However, the constant part of the objective function is irrelevant for the algorithm, and our implementations ignore it).

Now let us consider the edit constraints. Suppose the previous edit value for v was α , and the new edit value for v is β . The current tableau contains the

information that

$$v = \alpha + \delta_v^+ - \delta_v^-$$

and again we need to modify this, so that, instead,

$$v = \beta + \delta_v^+ - \delta_v^-$$

To do so we must replace every occurrence of

$$\delta_v^+ - \delta_v^-$$

by

$$\beta - \alpha + \delta_v^+ - \delta_v^-$$

taking proper account of the coefficients of δ_v^+ and δ_v^- . (Again, remember that δ_v^+ and δ_v^- come in pairs).

If either δ_v^+ or δ_v^- is basic, this simply involves appropriately modifying the equation in which they are basic. Otherwise, if both are nonbasic, then we need to change every equation of the form

$$x_i = c_i + a'_v \delta_v^+ - a'_v \delta_v^- + e$$

to

$$x_i = c_i + a'_v(\beta - \alpha) + a'_v \delta_v^+ - a'_v \delta_v^- + e.$$

Hence, modifying the tableau to reflect the new values of edit and stay constraints involves only changing the constant values in some equations. The modifications for stay constraints always result in a tableau in basic-feasible solved form, since it never makes a constant become negative. In contrast, the modifications for edit constraints may not.

To return to our example, suppose we pick up x_m with the mouse and move it to 60. Then we have that $\alpha = 50$ and $\beta = 60$, so we need to add 10 times the coefficient of $\delta_{x_m}^+$ to the constant part of every row. The modified tableau, after the updates for both the stays and edits, is

minimize $[0, 20] + [1, 2]\delta_{x_m}^+ + [1, -2]\delta_{x_m}^- + [0, 2]\delta_{x_l}^- + [0, 2]\delta_{x_r}^-$ subject to

$x_m = 60$	$+ \delta_{x_m}^+$	$- \delta_{x_m}^-$		
$x_l = 30$		$+ \delta_{x_l}^+$	$- \delta_{x_l}^-$	
$x_r = 90$	$+ 2\delta_{x_m}^+$	$- 2\delta_{x_m}^-$	$- \delta_{x_l}^+$	$+ \delta_{x_l}^-$
$s_1 = 50$	$+ 2\delta_{x_m}^+$	$- 2\delta_{x_m}^-$	$- 2\delta_{x_l}^+$	$+ 2\delta_{x_l}^-$
$s_3 = 10$	$- 2\delta_{x_m}^+$	$+ 2\delta_{x_m}^-$	$+ \delta_{x_l}^+$	$- \delta_{x_l}^-$
$\delta_{x_r}^+ = 20$	$+ 2\delta_{x_m}^+$	$- 2\delta_{x_m}^-$	$- \delta_{x_l}^+$	$+ \delta_{x_l}^- + \delta_{x_r}^-$
$s_2 = 40$			$+ \delta_{x_l}^+$	$- \delta_{x_l}^-$

Clearly, it is feasible and already in optimal form, and so we have incrementally resolved the problem by simply modifying constants in the tableaux. The new tableaux give the solution $\{x_m \mapsto 60, x_l \mapsto 30, x_r \mapsto 90\}$. So sliding the midpoint rightwards has caused the right point to slide rightwards as well, but twice as far. The resulting diagram is shown at the top of Figure 6.

Suppose we now move x_m from 60 to 90. The modified tableau is

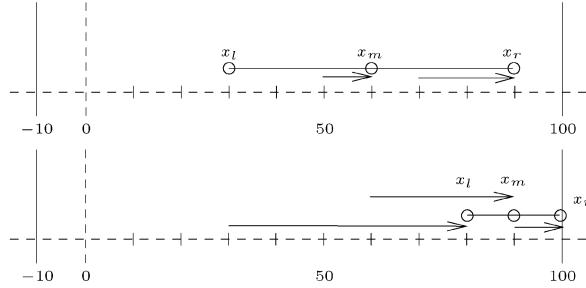


Fig. 6. Resolving the constraints.

minimize $[0, 60] + [1, 2]\delta_{x_m}^+ + [1, -2]\delta_{x_m}^- + [0, 2]\delta_{x_l}^- + [0, 2]\delta_{x_r}^-$ subject to

$$\begin{array}{rcl} x_m & = & 90 + \delta_{x_m}^+ - \delta_{x_m}^- \\ x_l & = & 30 + \delta_{x_l}^+ - \delta_{x_l}^- \\ x_r & = & 150 + 2\delta_{x_m}^+ - 2\delta_{x_m}^- - \delta_{x_l}^+ + \delta_{x_l}^- \\ \hline s_1 & = & 110 + 2\delta_{x_m}^+ - 2\delta_{x_m}^- - 2\delta_{x_l}^+ + 2\delta_{x_l}^- \\ s_3 & = & -50 - 2\delta_{x_m}^+ + 2\delta_{x_m}^- + \delta_{x_l}^+ - \delta_{x_l}^- \\ \delta_{x_r}^+ & = & 60 + 2\delta_{x_m}^+ - 2\delta_{x_m}^- - \delta_{x_l}^+ + \delta_{x_l}^- + \delta_{x_r}^- \\ s_2 & = & 40 + \delta_{x_l}^+ - \delta_{x_l}^- \end{array}$$

The tableau is no longer in basic-feasible-solved form, since the constant of the row for s_3 is negative, even though s_3 is supposed to be nonnegative. (In this solution, $x_r = 150$, so that the right endpoint has crashed through the $x_r \leq 100$ barrier).

Thus, in general, after updating the constants for the edit constraints, the simplex tableau C_S may no longer be in basic-feasible-solved form, since some of the constants may be negative. However, the tableau is still in basic form, so we can still read a solution directly from it as before. Also, because no coefficient has changed (in particular, the optimization function is the same), the resulting tableau reflects an optimal, but not feasible, solution.

We need to find a feasible and optimal solution. We could do so by adding artificial variables (as when adding a constraint, see Section 2.5), optimizing the sum of the artificial variables to find an initial feasible solution, and then reoptimizing the original problem.

But we can do much better. The process of moving from an *optimal and infeasible* solution to an *optimal and feasible* solution is exactly the dual of normal simplex algorithm, where we progress from a feasible and nonoptimal solution to feasible and optimal solution. Hence, we can use the *dual simplex algorithm* to find a feasible solution while staying optimal.

Solving the dual optimization problem starts from an infeasible-optimal-solved form tableau

minimize $e + \sum_{j=1}^m d_j y_j$ subject to

$$\bigwedge_{i=1}^n x_i = c_i + \sum_{j=i}^m a_{ij} y_j$$

where some c_i may be negative for rows with nonnegative basic variables (accounting for the tableau's infeasibility) and each d_j is nonnegative (so it is optimal).

```

re_optimize( $C_S, f$ )
foreach stay :  $v \in C$ 
    if  $\delta_v^+$  or  $\delta_v^-$  is basic in row  $i$  then  $c_i := 0$  endif
endfor
foreach edit :  $v \in C$ 
    let  $\alpha$  and  $\beta$  be the previous and current edit values for  $v$ 
    let  $\delta_v^+$  be  $y_j$ 
    foreach  $i \in \{1, \dots, n\}$ 
         $c_i := c_i + a_{ij}(\beta - \alpha)$ 
    endfor
endfor
repeat
    % Choose variable  $x_I$  to become non-basic
    choose  $I$  where  $c_I < 0$ 
    if there is no such  $I$ 
        return true
    endif
    % Choose variable  $y_J$  to become basic
    if for each  $j \in \{1, \dots, m\}$   $a_{Ij} \leq 0$  then
        return false
    endif
    choose  $J \in \{1, \dots, m\}$  such that
         $d_J/a_{IJ} = \min_{j \in \{1, \dots, m\}} \{d_j/a_{Ij} \mid a_{Ij} > 0\}$ 
     $l := (x_I - c_I - \sum_{j=1, j \neq J}^m a_{Ij}y_j)/a_{IJ}$ 
    replace  $y_J$  by  $l$  in  $f$ 
    for each  $i \in \{1, \dots, n\}$ 
        if  $i \neq I$  then replace  $y_J$  by  $l$  in row  $i$  endif
    endfor
    replace the  $I^{th}$  row by  $y_J = l$ 
until false

```

Fig. 7. Dual simplex reoptimization of constraints C_S with objective function f in infeasible-optimal-solved form.

The dual simplex algorithm selects an exit variable by finding a row I with nonnegative basic variable x_I and negative constant c_I . The entry variable is the variable y_J such that the ratio d_J/a_{IJ} is the minimum of all d_j/a_{Ij} where a_{Ij} is positive.³ This selection criteria ensures that when pivoting we stay at an optimal solution. The pivot replaces y_j by

$$-1/a_{Ij}(-x_I + c_I + \sum_{j=1, j \neq J}^m a_{Ij}y_j),$$

and is performed as in the (primal) simplex algorithm. The algorithm is shown in Figure 7.

Continuing the example above, we select the exit variable s_3 —the only non-negative basic variable for a row with negative constant. We find that $\delta_{x_l}^+$ has the minimum ratio since its coefficient in the optimization function is 0, so it will be the entry variable. Replacing $\delta_{x_l}^+$ everywhere by $50 + s_3 + 2\delta_{x_m}^+ - 2\delta_{x_m}^- + \delta_{x_l}^-$, we obtain the tableau:

³Recall that the d_j will be lexicographic coefficients, so this minimization is with respect to the lexicographic order.

minimize $[0, 60] + [1, 2]\delta_{x_m}^+ + [1, -2]\delta_{x_m}^- + [0, 2]\delta_{x_l}^- + [0, 2]\delta_{x_r}^-$ subject to

$$\begin{array}{rccc} x_m & = & 90 & +\delta_{x_m}^+ & -\delta_{x_m}^- \\ x_l & = & 80 + s_3 & +2\delta_{x_m}^+ & -2\delta_{x_m}^- \\ x_r & = & 100 - s_3 & & \\ \hline s_1 & = & 10 - 2s_3 & -2\delta_{x_m}^+ & +2\delta_{x_m}^- \\ \delta_{x_l}^+ & = & 50 + s_3 & +2\delta_{x_m}^+ & -2\delta_{x_m}^- + \delta_{x_l}^- \\ \delta_{x_r}^+ & = & 10 - s_3 & & +\delta_{x_r}^- \\ s_2 & = & 90 + s_3 & +2\delta_{x_m}^+ & -2\delta_{x_m}^- \end{array}$$

The tableau is feasible (and of course still optimal) and represents the solution $\{x_m \mapsto 90, x_r \mapsto 100, x_l \mapsto 80\}$. So by sliding the midpoint further right, the rightmost point hits the wall and the left point slides right to satisfy the constraints. The resulting diagram is shown at the bottom of Figure 6.

To summarize, incrementally finding a new solution for new input variables involves updating the constants in the tableaux to reflect the updated stay constraints, then updating the constants to reflect the updated edit constraints, and finally reoptimizing if needed. In an interactive graphical application, this dual optimization method typically requires a pivot only when one part of the figure first hits or first moves away from a barrier. The intuition behind this is that when a constraint first becomes unsatisfied, the value of one of its error variables will become nonzero, and hence the variable will have to enter the basis; conversely, when a constraint first becomes satisfied, we can move one of its error variables out of the basis.

In the example, pivoting occurred when the right point x_r came up against a barrier. Thus, if we picked up the midpoint x_m with the mouse and smoothly slid it rightwards, 1 pixel every screen refresh, only one pivot would be required in moving from 50 to 95. This behavior is why the dual optimization is well suited to this problem and leads to efficiently resolving the hierarchical constraints.

2.5 Incrementality: Adding a Constraint

We now describe how to add a new constraint incrementally. This technique is also used in our implementation to find an initial basic-feasible-solved form for the original simplex problem, by starting from an empty constraint set and adding the constraints one at a time.

As an example, suppose we wish to require that the midpoint be centered, i.e., to add a required constraint $x_m = 50$ to the final tableau given in Section 2.2, page 12. For reference, that tableau is

$$\begin{array}{rcl} x_l & = & 90 - s_1 - s_3 \\ x_m & = & 95 - \frac{1}{2}s_1 - s_3 \\ x_r & = & 100 - s_3 \\ \hline s_2 & = & 100 - s_1 - s_3 \end{array}$$

If we substitute for each of the basic variables in $x_m = 50$ (namely x_m), we obtain the equation $45 - \frac{1}{2}s_1 - s_3 = 0$. In order to add this constraint straightforwardly to the tableau, we create a new nonnegative variable a called an *artificial variable*. (This technique is simply an incremental version of the

operation in phase I of the two-phase simplex algorithm). We let $a = 45 - \frac{1}{2}s_1 - s_3$ be added to the tableau (clearly, this gives a tableau in basic-feasible-solved form) and then minimize the value of a . If a takes the value 0, then we have obtained a solution to the problem with the added constraint, and we can then eliminate the artificial variable altogether, since it is a parameter (and hence takes the value 0). This is the case for our example; the resulting tableau is

$$\begin{array}{r} x_l = 0 + s_3 \\ x_m = 50 \\ x_r = 100 - s_3 \\ \hline s_1 = 90 - 2s_3 \\ s_2 = 10 + s_3 \end{array}$$

In general, to add a new constraint to the tableau, if it is an inequality we first convert it to an equation by adding a slack variable. Next, we use the current tableau to substitute out all the basic variables. This gives an equation $l = 0$, where l is a linear expression. If the constant part of l is negative, we multiply both sides of the equation $l = 0$ by -1 , so that the constant becomes nonnegative. If l contains an unrestricted variable, we use it to substitute for that variable and add the equation to the tableau above the line (i.e., to C_U). Otherwise, we create a restricted artificial variable a , add the equation $a = l$ to the tableau below the line (i.e., to C_S), and minimize a . If the resulting minimum is not zero, then the constraints are unsatisfiable. Otherwise a is either parametric or basic. If a is parametric, the column for it can be simply removed from the tableau. If it is basic, the row must have constant 0 (since we were able to achieve a value of 0 for our objective function, which is equal to a). If the row is just $a = 0$, it can be removed. Otherwise, $a = 0 + bx + \dots$, where $b \neq 0$. We can then pivot x into the basis using this row and remove the column for a .

If the equation being added contains any unrestricted variables after substituting out all the basic variables, as described above, we do not need to use an artificial variable. Not only that, we *could not* use an artificial variable, since we cannot put an equation in C_S that contains an unrestricted variable. In some other cases we can avoid using an artificial variable for efficiency, even though it would be permissible to use one. We can avoid using an artificial variable if we can choose a subject for the equation from among its current variables.

Here are the rules for choosing a subject. (These are to be used after replacing any basic variables with their defining expressions). We are adding the constraint $l = 0$ to the tableau. If necessary, normalize l by multiplying by -1 so that its constant part is nonnegative. We want to pick a variable in l to be the subject of an equation, so that we can add the row $v = l'$, where l' the result of solving $l = 0$ for v .

- If l contains any unrestricted variables, we must choose an unrestricted variable as the subject.
- If the subject is new to the solver, we will not have to do any substitutions, so we prefer new variables to ones that are currently noted as parametric.
- If l contains only restricted variables, if there is a variable v in l that has

a negative coefficient and that is new to the solver, we can pick v as the subject. (Adding the equation $v = l'$ that is the result of solving $l = 0$ for v will result in a basic feasible solved form, since the constant part of l' will be nonnegative).

—Otherwise use an artificial variable.

A consequence of these rules is that we can always add a nonrequired constraint to the tableau without using an artificial variable, because the equation will contain a positive and a negative error or slack variable, both of which are new to the solver, and which occur with opposite signs. (Constraints that are originally equations will have a positive and a negative error variable, while constraints that are originally inequalities will have one error variable and one slack variable, with opposite signs). This observation is good news for performance, since adding a nonrequired edit constraint is a common operation.

If the subject variable is a new error variable, or an artificial variable, we must reoptimize the resulting system after the addition. Also, if the subject is an artificial variable, we must remove it from the tableau C_S before reoptimizing (see Marriott and Stuckey [1998, p. 70] for details).

2.6 Incrementality: Removing a Constraint

We also want a method for incrementally removing a constraint from the tableaux. After a series of pivots have been performed, the information represented by the constraint may not be contained in a single row, so we need a way to identify the constraint's influence in the tableaux. To do this, we use a “marker” variable that is originally present only in the equation representing the constraint. We can then identify the constraint's effect on the tableaux by looking for occurrences of that marker variable. For inequality constraints, the slack variable s that we added to make it an equality serves as the marker (since s will originally occur only in that equation). For nonrequired equality constraints, either of its two error variables can serve as a marker—see Section 2.3. For required equality constraints, we add a “dummy” restricted variable to the original equation to serve as a marker, which we never allow to enter the basis (so that it always has value 0). In our running example, then, to allow the constraint $2x_m = x_l + x_r$ to be deleted incrementally we would have added a dummy variable d , resulting in $2x_m = x_l + x_r + d$. The simplex optimization routine checks for these dummy variables in choosing an entry variable and does not allow one to be selected. These dummy variables must be restricted, not unrestricted, because we might need to have some of them in the equations for restricted basic variables. (We did not include the variable d in the tableaux presented earlier to simplify the presentation).

Consider removing the constraint that x_l is 10 to the left of x_r . The slack variable s_1 , which we added to the inequality to make it an equation, records exactly how this equation has been used to modify the tableau. We can remove the inequality by pivoting the tableau until s_1 is basic and then simply drop the row in which it is basic.

In the tableau in Section 2.5 (obtained after adding the required constraint $x_m = 50$), s_1 is already basic, so removing it simply means dropping the row in

```

inc.addition( $C_U, C_S, C_I, f, s, c$ )
 $C_I^{old} := C_I$ 
if  $c$  is of the form  $l \geq 0$ 
  if  $s$  is required
     $c := l = s'$  where  $s'$  is a new slack variable
     $C_I := C_I \cup \{s' \geq 0\}$ 
  else
     $c := l = s' - \delta_c$  where  $\delta_c$  is a new error variable
     $C_I := C_I \cup \{s' \geq 0, \delta_c \geq 0\}$ 
     $f := f + s\delta_c$ 
  endif
else
  let  $c$  be of the form  $l = 0$  where the constant in  $l$  is  $\geq 0$ 
  if  $s$  is not required
     $c := l = \delta_c^+ - \delta_c^-$ 
     $C_I := C_I \cup \{\delta_c^+ \geq 0, \delta_c^- \geq 0\}$ 
     $f := f + s\delta_c^+ + s\delta_c^-$ 
  endif
endif
for each  $x = l$  in  $C_U \cup C_S$ 
  replace  $x$  in  $c$  by  $l$ 
endfor
if exists  $y \in vars(c) - vars(C_I)$ 
  let  $c$  be of the form  $y = l'$ 
  replace  $y$  by  $l'$  everywhere in  $C_U$ 
   $C_U := C_U \cup \{y = l'\}$ 
elseif exists  $y \in vars(c) - vars(C_I^{old})$ 
  where  $c$  is of the form  $y = l'$  where the constant in  $l'$  is  $\geq 0$ 
   $C_S := C_S \cup \{y = l'\}$ 
  replace  $y$  by  $l'$  in  $f$ 
   $(C_S, f) := simplex\_opt(C_S, f)$ 
else %%%  $c$  must be a required constraint
  let  $c$  be of form  $l' = 0$  where the constant in  $l'$  is  $\geq 0$ 
   $C_S := C_S \cup \{a = l'\}$  where  $a$  is a new artificial variable
   $(C_S, l'') := simplex\_opt(C_S, l'')$ 
  if constant part of  $l''$  is not zero then exception "unsatisfiable constraints"
    remove variable  $a$  from  $C_S$  (possibly pivoting)
     $(C_S, f) := simplex\_opt(C_S, f)$ 
  endif
endif
return  $(C_U, C_S, C_I, f)$ 

```

Fig. 8. Incremental addition and reoptimization of a constraint c with strength s to the system $C_U \cup C_S \cup C_I$ with objective function f .

which it is basic, obtaining

$$\begin{array}{rcl}
 x_l &=& 0 + s_3 \\
 x_m &=& 50 \\
 x_r &=& 100 - s_3 \\
 \hline
 s_2 &=& 10 + s_3
 \end{array}$$

If we wanted to remove this constraint from the tableau before adding $x_m = 50$ (i.e., the final tableau in Section 2.2), s_1 is a parameter. We make s_1 basic by treating it as an entry variable and (as usual) determining the most restrictive equation, then using that to pivot s_1 into the basis before finally removing the row.

There is such a restrictive equation in this example. However, if the marker variable does not occur in C_S , or if its coefficients in C_S are all nonnegative, then no equation restricts the value of the marker variable. If the marker variable does occur in one or more equations in C_S , always with a positive coefficient, pick the equation with the smallest ratio of the constant to the marker variable's coefficient. (The row with the marker variable will become infeasible after the pivot, but all the other rows will still be feasible, and we will be dropping the row with the marker variable. In effect, we are removing the nonnegativity restriction on the marker variable). Finally, if the marker variable occurs only in equations for unrestricted variables, we can choose any equation in which it occurs.

In the final tableau in Section 2.2, the row $s_2 = 100 - s_1 - s_3$ is the most constraining equation. Pivoting to let s_1 enter the basis and then removing the row in which it is basic, we obtain

$$\begin{aligned} x_l &= -10 + s_2 \\ x_m &= 45 + \frac{1}{2}s_2 - \frac{1}{2}s_3 \\ x_r &= 100 - s_3 \end{aligned}$$

In the preceding example, the marker variable had a negative coefficient. Here is an example that only has positive coefficients. The original constraints are

$$\begin{aligned} x &\geq 10 \\ x &\geq 20 \\ x &\geq 30 \end{aligned}$$

In basic-feasible-solved-form, this is

$$\begin{aligned} x &= 30 + s_3 \\ s_1 &= 20 + s_3 \\ s_2 &= 10 + s_3 \end{aligned}$$

where s_1 , s_2 , and s_3 are the marker (and slack) variables for $x \geq 10$, $x \geq 20$, and $x \geq 30$, respectively. This gives a solution for x of $x = 30$, which of course satisfies all of the original inequalities.

Suppose we want to remove the $x \geq 30$ constraint. We need to pivot to make s_3 basic. The equation that gives the smallest ratio is $s_2 = 10 + s_3$, so the entry variable is s_3 and the exit variable is s_2 , giving

$$\begin{aligned} x &= 20 + s_2 \\ s_1 &= 10 + s_2 \\ s_3 &= -10 + s_2 \end{aligned}$$

This tableau is now infeasible, but we drop the row with s_3 , giving

$$\begin{aligned} x &= 20 + s_2 \\ s_1 &= 10 + s_2 \end{aligned}$$

which is of course feasible.

A beneficial result of using marker variables is that redundant constraints can be represented and manipulated. Consider

$$\begin{aligned} x &\geq 10 \\ x &\leq 10 \end{aligned}$$

When converted to basic-feasible-solved form, each $x \geq 10$ constraint gets a separate slack variable, which is used as the marker variable for that constraint.

$$\begin{array}{r} x = 10 + s_1 \\ s_2 = 0 + s_1 \end{array}$$

To delete the second $x \geq 10$ constraint, we simply drop the $s_2 = 0 + s_1$ row. To delete the first $x \geq 10$ constraint, we pivot, making s_1 basic and s_2 parametric:

$$\begin{array}{r} x = 10 + s_2 \\ s_1 = 0 + s_2 \end{array}$$

and then drop the $s_1 = 0 + s_2$ row.

A consequence of directly representing redundant constraints is that they must all be removed to eliminate their effect. (This seems to be a more desirable behavior for the solver than removing redundant constraints automatically; although if the latter were desired, the solver could be modified to do this.) Another consequence is that when adding a new constraint, we would never decide that it was redundant and not add it to the tableau.

Before we remove the constraint, there may be some stay constraints that were previously unsatisfied. If we just removed the constraint, these could come into play, so instead we reset all of the stays so that all variables are constrained to stay at their current values.

Also, if the constraint being removed is not required, we need to remove the error variables for it from the objective function. To do this, we add the following to the expression for the objective function:

$$-1 \times l \times s \times w$$

where l is the error variable if it is parametric, or else l is its defining expression if it is basic, s is the unit symbolic weight for the constraint's strength, and w is its weight at the given strength.

If we allow nonrequired constraints other than stays and edits, we must also reoptimize after deleting a constraint, since a nonrequired constraint might have become satisfiable (or more nearly satisfiable). The complete algorithm is shown in Figure 9.

2.7 Strict Inequalities

Cassowary supports linear equations and nonstrict inequalities, but not strict inequalities. It is straightforward to add support for strict inequalities, but doing this would have some consequences that users might not expect, and so we believe it is not a good idea. The issue is that if strict inequalities were allowed, it would be easy to construct constraint hierarchies such that there were no solutions to the hierarchy, even though the required constraints could be satisfied. For example, consider the constraint hierarchy {required $x > 0$, weak $x = 0$ }. For any candidate solution $\{x \mapsto \epsilon\}$, the solution $\{x \mapsto \epsilon/2\}$ more nearly satisfies the weak constraint—and so there is no solution to the hierarchy. (An analogous problem occurs with the simplex algorithm itself—consider minimizing x subject to $x > 0$).

```

inc_deletion( $C_U, C_S, C_I, f, v$ )
foreach stay :  $c \in C$ 
    if  $\delta_c^+$  or  $\delta_c^-$  is basic in row  $i$  then  $c_i := 0$  endif
endfor
if  $v$  is parametric variable  $y_J$ 
    if  $y_J \in vars(C_S)$ 
        choose  $I \in \{1, \dots, n\}$  such that
             $-c_I/a_{IJ} = \min_{i \in \{1, \dots, n\}} \{-c_i/a_{iJ} \mid a_{iJ} < 0\}$ 
        remove row  $I$  from  $C_S$ 
    else
        choose arbitrary row  $y_J = l$  in  $C_U$ 
        remove the row from  $C_U$ 
        replace  $y_J$  by  $l$  in  $f$ 
        for each row in  $C_U \cup C_S$ 
            replace  $y_J$  by  $l$  in row  $i$  endif
        endfor
    endif
endif
 $\Delta :=$  error variables of constraint  $c$  for which  $v$  is marker
 $f := f$  with coefficients of  $\Delta$  variables set to 0
 $(C_S, f) :=$  simplex_opt( $C_S, f$ )
return  $(C_U, C_S, C_I, f)$ 

```

Fig. 9. Incremental deletion and reoptimization for a constraint with marker variable v deleted from $C_U \cup C_S \cup C_I$ with objective function f .

If strict inequalities were nevertheless desired, they could be supported in the same manner as described by either Jaffar et al. [1992] or Van Hentenryck and Graf [1990]. The algorithm would additionally check for error variables in the objective function with infinitesimal values; and if one were found, raise a “no solution to constraints” exception.

2.8 Algorithm Summary

The Cassowary algorithm supports three basic operations:

- incremental addition of a new constraint, and reoptimization to find a weighted-sum-better solution;
- incremental deletion of a previous constraint (using its marker variable), and reoptimization to find a weighted-sum-better solution; and
- incremental resolving to find a weighted-sum-better solution after modification of edit constraints, including special treatment of stay constraints may change after every resolve.

The third operation is the most crucial in terms of performance, since it occurs in realtime as the user moves the mouse.

While we have used only three constraint strengths in the examples in this article, both the algorithm and implementation are general, and can be used with an arbitrary number of strengths.

3. EMPIRICAL EVALUATION

Cassowary has been implemented in Smalltalk, C++, and Java. In this section we first describe timings for a large, real application (Scwm, the Scheme Constraints Window Manager), and then a set of benchmarks, including two graphical examples and some comparisons between Cassowary and a local propagation solver.

The C++ and Java tests were performed on an 800 MHz Pentium machine running RedHat/Linux 7.0, using G++-2.96-69 and the Sun HotSpot JVM 1.3.0, respectively. All the Smalltalk tests were run using OTI Smalltalk R4.0, on an 800 MHz Pentium machine running Windows 2000. (Cassowary has also been ported to the Squeak dialect of Smalltalk).

3.1 Scwm Performance

For these tests, Scwm ran on a 2560x1024 Xinerma-enabled XFree86 4.0 display. Scwm is written in C and Guile/Scheme and embeds the C++ implementation of Cassowary using a hand-written Guile/Scheme wrapper instrumented to support recording the time spent inside the various constraint operations.

Adding 25 new XTerm windows and their implicit stay constraints required 3.6 ms per window to add the 8 initial constraints, and 0.14 ms for each of 72 resolves. Moving arbitrary windows around the desktop without any constraints relating the windows required about 0.48 ms for each resolve.

We then made a 2×2 window tiling of four windows and constrained a fifth window to the left of that tiled collection and a sixth to the right of it. Adding each of those constraints took 0.9 ms on average. Subsequently, we moved those six related windows around the screen a bit and measured that each resolve took only 0.66 ms.

3.2 Graphical Benchmarks

Both of the examples in this section were implemented in Smalltalk.

For the bounded quadrilateral demonstration shown in Figure 1, it took an average of 14.97 msec to resolve the constraints and refresh the display for a new mouse position. There were an average of 0.2 pivots per refresh (the mouse was being moved quickly, so that there were a relatively large number of collisions with the sides of the windows). Of the 14.97 msec, 1.4% was spent in the constraint solver—in other words, almost all of the time was taken by graphics refresh and input handling.

Figure 10 shows a binary tree of height 7, which exercises the constraint solver more vigorously. Each node in the tree has a weak stay, and is required to lie within the confines of the window. Also, the y values of the two children of each parent are required to be equal, and to be at least 10 pixels below the parent's y value. The x value of the parent node is required to be halfway between the x values of its children. There are a total of 1015 constraints for the height 7 tree. When the user moved the root node of the tree, it took an average of 41.4 msec to resolve the constraints and refresh the display for a new mouse position, of which 33% was spent in the constraint solver. There were an average of 5.9 pivots per refresh (so many more collisions than

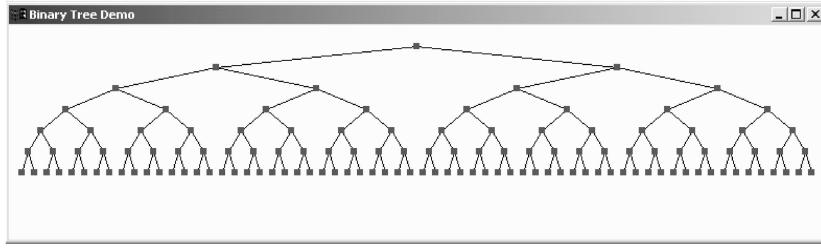


Fig. 10. A binary tree of height 7.

with the quadrilateral—as with the quadrilateral, the mouse was being moved very quickly).

3.3 Local Propagation Comparisons

Efficient incremental algorithms have been developed that solve systems of constraints using multiway local propagation, e.g., DeltaBlue [Sannella et al. 1993], SkyBlue [Sannella 1994], and QuickPlan [Vander Zanden 1996]. In this section, we compare the performance of DeltaBlue and Cassowary on the same constraint problems. DeltaBlue is much more efficient—but it only handles equality constraints, not inequalities, and requires that the constraint graph be acyclic. We report the results for three benchmarks: the chain, the tree, and the star. These have all been used in previously published descriptions of local propagation algorithms. In each case the times are for Smalltalk implementations of DeltaBlue and Cassowary, both running on the same 800 MHz machine, to allow a head-to-head comparison.

3.3.1 Chain Benchmark. In the chain benchmark, we construct a long chain of required equality constraints $x_1 = x_2 = \dots = x_{1000}$. There is a weak stay on x_{1000} . In DeltaBlue, there is implicitly also a weakest stay on every variable; in Cassowary, the other variables have no additional constraints. In DeltaBlue, the equality constraints are satisfied using the methods $x_{999} \leftarrow x_{1000}, x_{998} \leftarrow x_{999}, \dots, x_1 \leftarrow x_2$; in other words, the values of x_1, x_2, \dots, x_{999} are all ultimately determined by x_{1000} . In Cassowary, the error variables $\delta_{x_{1000}}^+$ and $\delta_{x_{1000}}^-$ for the weak stay on x_{1000} are parametric, and x_1, \dots, x_{1000} are basic.

We then add a strong edit constraint to x_1 . In DeltaBlue, this causes each of the 999 equality constraints to flip and select the other satisfaction method: $x_2 \leftarrow x_1, x_3 \leftarrow x_2, \dots, x_{1000} \leftarrow x_{999}$. In Cassowary, the error variables $\delta_{x_1}^+$ and $\delta_{x_1}^-$ for the edit constraint become basic, and one of the error variables for the stay becomes parametric—so that each row for x_1, \dots, x_{1000} must be changed. Thus, adding the edit constraint is a pessimal case for both solvers, in that every constraint must be touched. We then repeatedly change the value of x_1 using the edit constraint. The results are shown in Figure 11.

3.3.2 Star Benchmark. In the Star benchmark, there are n input variables x_i , n output variables y_i , and a single offset z . These variables are related by n

Action	DeltaBlue	Cassowary
build chain with 1000 variables (total time)	70.6	38004.0
add strong edit constraint	5.4	61.0
initial solution	4.5	12528.0
change value of x_1 one time	0.4	5.9

Fig. 11. Chain benchmark for chain of length 1000; all times are in msec.

Action	DeltaBlue	Cassowary
build star with $n = 100$ (total time)	26.3	1483.0
add strong edit constraint to offset	0.7	< 1.0
initial solution	0.6	2003.0
change value of offset one time	0.1	1.5

Fig. 12. Star benchmark for $n = 100$; all times are in msec.

Action	DeltaBlue	Cassowary
build tree of depth 10 (total time)	231.5	2413.0
add strong edit constraint to root	0.3	241.0
initial solution	0.1	891.0
change value of root one time	0.01	6.4

Fig. 13. Tree benchmark for $n = 10$; all times are in msec.

required constraints $x_i + z = y_i$. Each input variable x_i has a medium stay, and each output variable y_i has a weak stay. We then add a strong edit constraint to the offset. (In DeltaBlue, this forces the solver to choose a different method for each of the required offset constraints). Then, as the value of the offset is edited, each output variable changes as well to satisfy the constraints. The results are shown in Figure 12. (In the original benchmark, each input variable is multiplied by a scaling factor, rather than being added to an offset—but these constraints would be nonlinear. DeltaBlue’s performance is not affected by this change).

3.3.3 Tree Benchmark. In the Tree benchmark, a binary tree of depth n is constructed. The value at each interior node is constrained to be the sum of the values of its children. There is also a weak stay constraint on each leaf. We then add a strong edit constraint to the root. This benchmark is a particularly favorable case for DeltaBlue, but not for Cassowary: In DeltaBlue, adding the strong edit constraint requires touching only $\log n$ of the constraints, while in Cassowary it requires only one pivot, but on a constraint of size n . The results are shown in Figure 13.

3.4 Remarks

In all of our interactive graphical applications, during manipulation the time to refresh the picture and handle input dominated the constraint satisfaction time. Interaction times remain satisfactory for diagrams with 1000 constraints. Cassowary is much slower than a local propagation solver when run on the same problem—an order of magnitude for resolve time, and slower still for finding the initial solution. In many situations this will not matter, since the constraint satisfaction time in either case will be a small portion of the overall time. In

situations where this is an issue, however, these results indicate that a hybrid solver architecture would be beneficial, in which there are cooperating subsolvers, including one to handle local propagation constraints and another to handle linear equalities and inequalities. Then, when possible, constraints would be allocated to the more efficient local propagation solver. A cooperating-subsolver architecture is used in Ultraviolet, but some adaptations would be needed to use it efficiently with Cassowary [Borning and Freeman-Benson 1998].

4. APPLICATIONS

The various implementations of Cassowary are in active use. A Scheme wrapping of the C++ implementation is used in the Scheme Constraints Window Manager. We also embedded the C++ implementation in a prototype Web browser that supports Constraint Cascading Style Sheets (CCSS) [Badros et al. 1999], our constraint-based extension to Cascading Style Sheets. CCSS allows both the designer and the viewer to place constraints on the appearance of a Web page—for example, on page width or font size. Some of these constraints can be requirements, and others preferences. As a result, the final appearance of the page is, in effect, the result of an arbitration between the desires of the designer and the viewer, where this arbitration is carried out by solving the combined sets of constraints. An earlier, related system is described in Borning et al. [1997; 2000]. Just as CCSS extends Cascading Style Sheets with constraints, so CSVG, the Constraint Structured Vector Graphics standard [Badros et al. 2001], extends SVG with constraints. A SVG viewer, augmented with Cassowary, allows diagrams to be laid out in different ways, depending on the constraints from the author and the viewer. In other work, a demonstration, Constraint Drawing Application of Cassowary, written by Michael Noth using Java implementation, is included with the Cassowary toolkit. Cassowary has also been used in a noninteractive application to perform consistency checks in a planning application [Wolfman and Weld 1999].

5. CONCLUSION

The Cassowary algorithm is an efficient constraint solver for interactive user interface applications, which handles simultaneous linear equations and inequalities. Due to the minimal update of the tableau that is performed, it is (perhaps surprisingly) fast on the operation of incrementally resolving the system. That operation's efficiency is crucial for interactive redrawing diagrams during editing.

Additionally, because Cassowary handles cycles in the constraint graph without difficulty, users of the Cassowary implementations can concentrate on exploiting the additional expressiveness that the library provides; the declarative nature of constraints is not undermined by a need to reformulate the problem to avoid cycles. Cassowary has proven to be efficient and expressive enough to be used in many applications.

APPENDIX

A. IMPLEMENTATION DETAILS

A.1 Principal Classes

The principal classes in our implementations are as follows. All the classes start with “CI” for “Constraint Library” and are, of course, direct or indirect subclasses of `Object` in the Smalltalk and Java implementations.

```

Object
  CIAbstractVariable
    CIDummyVariable
    CIOjectiveVariable
    CISlackVariable
    CIVariable
  CIConstraint
    CIEditOrStayConstraint
      CIEditConstraint
      CIStayConstraint
    CILinearConstraint
      CILinearEqualityConstraint
      CILinearInequalityConstraint
    CILinearExpression
    CITableau
      CISimplexSolver
    CIStrength
    CISymbolicWeight
  
```

Some of these classes make use of the *Dictionary* (or *map*) abstract data type: dictionaries have keys and values and permit efficiently finding the value for a given key, and adding or deleting key/value pairs. We can also iterate through all keys, all values, or all key/value pairs.

A.2 Solver Protocol

The solver itself is represented as an instance of `CISimplexSolver`. Its public message protocol is as follows.

`addConstraint(CIConstraint cn)`. Incrementally add the linear constraint `cn` to the tableau. The constraint object contains its strength.

`removeConstraint(CIConstraint cn)`. Remove the constraint `cn` from the tableau. Also remove any error variables associated with `cn` from the objective function.

`addEditVar(CIVariable v, CIStrength s)`. Add an edit constraint of strength `s` on variable `v` to the tableau so that `suggestValue` (see below) can be used on that variable after a `beginEdit()`.

`removeEditVar(CIVariable v)`. Remove the previously added edit constraint on variable `v`. The `endEdit` message automatically removes all the edit variables as part of terminating an edit manipulation.

`beginEdit()`. Prepare the tableau for new values to be given to the currently-edited variables. The `addEditVar` message should be used before calling `beginEdit`, and `suggestValue` and `resolve` should be used only after `beginEdit` has been invoked, but before the required matching `endEdit`.

`suggestValue(CIVariable v, double n)`. Specify a new desired value n for the variable v . Before this call, v needs to have been added as a variable of an edit constraint (either by `addConstraint` of a hand-built `EditConstraint` object, or more simply using `addEditVar`).

`endEdit()`. Denote the end of an edit manipulation, thus removing all edit constraints from the tableau. Each `beginEdit` call must be matched with a corresponding `endEdit` invocation and the calls may be nested properly.

`resolve()`. Try to resolve the tableau given the newly specified desired values. Calls to `resolve` should be sandwiched between a `beginEdit()` and an `endEdit()`, and should occur after new values for edit variables are set using `suggestValue`.

`addPointStays(Vector points)`. This method is a bit of a kludge, and addresses the desire to satisfy the stays on both the x and y components of a given point, rather than on the x component of one point and the y component of another. The argument `points` is an array of points, whose x and y components are constrainable variables. This method adds a weak stay constraint to the x and y variables of each point. The weights for the x and y components of a given point are the same. However, the weights for successive points are each smaller than those for the previous point (1/2 of the previous weight). The effect of this is to encourage the solver to satisfy the stays on both the x and y of a given point rather than the x stay on one point and the y stay on another; see Borning and Badros [2000] for more discussion of this issue and a proposal for a cleaner solution.

`setAutoSolve(boolean f)`. Choose whether the solver should automatically optimize and set external variable values after each `addConstraint` or `removeConstraint`. By default, auto-solving is on, but passing `false` to this method will turn it off (until later turned back on by passing `true` to this method). When auto-solving is off, `solve` (below) or `resolve` must be invoked to see changes to the `CIVariables` contained in the tableau.

`isAutoSolving() returns boolean`. Return `true` if and only if the solver is auto-solving, `false` otherwise.

`solve()`. Optimize the tableau and set the external `CIVariables` contained in the tableau to their new values. This method need only be invoked if auto-solving has been turned off. It never needs to be called after a `resolve` method invocation.

`reset()`. Reinitialize the solver from the original constraints, thus getting rid of any accumulated numerical problems. (It is not yet clear how often such problems arise, but we provide the method just in case.)

A.2.1 Variables. `CIAbstractVariable` and its subclasses represent various kinds of constrained variables. `CIAbstractVariable` is an abstract class; that is, it is just used as a superclass of other classes, we do not make instances of `CIAbstractVariable` itself. `CIAbstractVariable` defines the message protocol for constrainable variables. Its only instance variable is `name`, which is a string name

Class	isDummy	isExternal	isPivotal	isRestricted
CIDummyVariable	true	false	false	true
ClVariable	false	true	false	false
ClSlackVariable	false	false	true	true
ClObjectiveVariable	false	false	false	false

Fig. 14. Subclasses of ClAbstractVariable.

for the variable. (This field is used for debugging and constraint understanding tasks.)

Instances of the concrete ClVariable subclass of ClAbstractVariable are what the user of the solver sees (hence it was given a nicer class name). This class has an instance variable value that holds the value of this variable. Users of the solver can send one of these variables the message value to get its value.

The other subclasses of ClAbstractVariable are used only within the solver. They do not hold their own values—rather, the value is just given by the current tableau. None of them have any additional instance variables.

Instances of ClSlackVariable are restricted to be nonnegative. They are used as the slack variable when converting an inequality constraint to an equation and for the error variables to represent nonrequired constraints.

Instances of CIDummyVariable are used as a marker variable to allow required equality constraints to be deleted. (For inequalities or nonrequired constraints, the slack or error variable is used as the marker.) These dummy variables are never pivoted into the basis.

An instance of ClObjectiveVariable is used to index the objective row in the tableau. (Conventionally this variable is named z .) This kind of variable is just for convenience—the tableau is represented as a dictionary (with some additional cross-references). Each row is represented as an entry in the dictionary; the key is a basic variable and the value is an expression. So an instance of ClObjectiveVariable is the key for the objective row. The objective row is unique, in that the coefficients of its expression are CISymbolicWeights in the Smalltalk implementation, not just ordinary real numbers. However, the C++ and Java implementations convert CISymbolicWeights to real numbers to avoid dealing with ClLinearExpressions parameterized on the type of the coefficient; see Section A.2.5 for more details.

All variables understand the following messages: isDummy, isExternal, isPivotal, and isRestricted. They also understand messages to get and set the variable's name.

For isDummy, instances of CIDummyVariable return true and others return false. The solver uses this message to test for dummy variables. It will not choose a dummy variable as the subject for a new equation, unless all the variables in the equation are dummy variables. (The solver also will not pivot on dummy variables, but this is handled by the isPivotal message.)

For isExternal, instances of ClVariable return true and others return false. If a variable responds true to this message, it means that it is known outside the solver, and so the solver needs to give it a value after solving is complete.

For `isPivotable`, instances of `CISlackVariable` return true and others return false. The solver uses this message to decide whether it can pivot on a variable.

For `isRestricted`, instances of `CISlackVariable` and of `CIDummyVariable` return true, and instances of `CIVariable` and `CIOjectiveVariable` return false. Returning true means that this variable is restricted to being nonnegative.

A variable's significance is largely just its identity (as mentioned above, variables have little state: a name for debugging and a value for instances of `CIVariable`). The only other messages that variables understand are some messages to `CIVariable` for creating constraints; see Section A.2.4.

A.2.2 Linear Expressions. Instances of the class `CILinearExpression` hold a linear expression and are used in building and representing constraints and in representing the tableau. A linear expression holds a dictionary of variables and coefficients (the keys are variables and the values are the corresponding coefficients). Only variables with nonzero coefficients are included in the dictionary; if a variable is not in this dictionary its coefficient is assumed to be zero. The other instance variable is a constant. So to represent the linear expression $a_1x_1 + \dots + a_nx_n + c$, the dictionary would hold the key x_1 with value a_1 , etc., and the constant c .

Linear expressions understand a large number of messages. Some of these are for constraint creation (see Section A.2.4). The others are to substitute an expression for a variable in the constraint, to add an expression, to find the coefficient for a variable, and so forth.

A.2.3 Constraints. There is an abstract class `CIConstraint` that serves as the superclass for other concrete classes. It defines two instance variables: `strength` and `weight`. The variable `strength` is the strength of this constraint in the constraint hierarchy (and should be an instance of `CIStrength`), while `weight` is a float indicating the actual weight of the constraint at its indicated strength, or `nil/null` if it does not have a weight. (Weights are only relevant for weighted-sum-better comparators, not for locally-error-better ones.)

Constraints understand various messages that return true or false regarding some aspect of the constraint, such as `isRequired`, `isEditConstraint`, `isStayConstraint`, and `isInequality`.

`CILinearConstraint` is an abstract subclass of `CIConstraint`. It adds an instance variable `expression`, which holds an instance of `CILinearExpression`. It has two concrete subclasses. An instance of `CILinearEquation` represents the linear equality constraint expression = 0.

An instance of `CILinearInequality` represents the constraint expression ≥ 0 .

The other part of the constraint class hierarchy is for edit and stay constraints (both are represented explicitly). `CIEditOrStayConstraint` has an instance field `variable`, which is the `CIVariable` with the edit or stay. Otherwise all that the two concrete subclasses do is respond appropriately to the messages `isEditConstraint` and `isStayConstraint`.

This constraint hierarchy is also intended to allow extension to include local propagation constraints (which would be another subclass of `CIConstraint`);

otherwise we could have made everything a linear constraint, and eliminated the abstract class `CIConstraint` entirely.

A.2.4 Constraint Creation. This section describes a mechanism to allow constraints to be defined easily by programmers. The convenience afforded by our toolkit varies among languages. Smalltalk's dynamic nature makes it the most expressive. C++'s operator overloading still permits using natural infix notation. Java, however, requires using ordinary methods, and leaves us with the single option of prefix expressions when building constraints.

In Smalltalk, the messages `+`, `-`, `*`, and `/` are defined for `CIVariable` and `CILinearExpression` to allow convenient creation of constraints by programmers. Also, `CIVariable` and `CILinearExpression`, as well as `Number`, define `cnEqual:`, `cnGEQ:`, and `cnLEQ:` to return linear equality or inequality constraints. Thus, the Smalltalk expression `3*x + 5 cnLEQ: y` returns an instance of `CILinearEquality` representing the constraint $3x + 5 \leq y$. The expression is evaluated as follows: the number 3 gets the message `*x`. Since `x` is not a number, 3 sends the message `*3` to `x`. `x` is an instance of `CIVariable`, which understands `*` to return a new linear expression with a single term, namely itself times the argument. (If the argument is not a number, it raises an exception that the expression is nonlinear.) The linear expression representing $3x$ gets the message `+` with the argument 5, and returns a new linear expression representing $3x + 5$. This linear expression gets the message `cnLEQ:` with the argument `y`. It computes a new linear expression representing $y - 3x - 5$, and then returns an instance of `CILinearInequality` with this expression.

(It is tempting to make this nicer by using the `=`, `<=`, and `>=` messages, so that we could write `3*x + 5 <= y` instead, but because the rest of Smalltalk expects `=`, `<=`, and `>=` to perform a test and return a boolean, rather than to return a constraint, this would not be a good idea).

Similarly, in C++ the arithmetic operators are overloaded to build `CILinearExpressions` from `CIVariables` and other `CILinearExpressions`. Actual constraints are built using various constructors for `CILinearEquation` or `CILinearInequality`. An enumeration defines the symbolic constants `cnLEQ` and `cnGEQ` to approximate the Smalltalk interface. For example,

```
CILinearInequality cn(3*x + 5, cnLEQ, y); // C++
```

builds the constraint `cn` representing $3x + 5 \leq y$. In Java, the same constraint would be built as follows:

```
CILinearInequality cn =
    new CILinearInequality(CL.Plus(CL.Times(x,3),5), CL.LEQ, y);
```

Although the Java implementation makes it more difficult to express programmer-written constraints, this inconvenience is relatively unimportant when the solver is used in conjunction with graphical user interfaces for specifying the constraints.

A.2.5 Symbolic Weights and Strengths. The constraint hierarchy theory allows an arbitrary, finite number of strengths of constraint. In practice, however,

programmers use a small number of strengths in a stylized way. The implementation therefore includes a small number of predefined strengths, and for efficiency the maximum number of strengths is defined as a compile-time constant. (This constant can be changed easily, but we have not found it necessary to do so in practice.)

The strengths provided in the current release are as follows:

required. Required constraints must be satisfied and should be used only for relationships that make no sense unless they are exactly met. The most common use of a required constraint is to give a shorthand name to an expression such as `win.right = win.left + win.width`

strong. This strength is conventionally used for edit constraints.

medium. This strength can be used for strong stays constraints; for example, we might put medium-strength stay constraints on the width and height of an object, and weak stay constraints on its position to represent our preference that the object move instead of change size when either is possible to maintain the stronger constraints.

weak. This strength is used for stay constraints.

Each strength category is represented as an instance of `CIStrong`.

A related class is `CISymbolicWeight`. As mentioned in Section 2.3, the objective function is formed as the weighted sum of the positive and negative errors for the nonrequired constraints. The weights should be such that the stronger constraints totally dominate the weaker ones. In general, to pick a real number for the weight, we need to know how big the values of the variables can be. To avoid this problem altogether, in the Smalltalk and C++ implementations we use symbolic weights and a lexicographic ordering for the weights rather than real numbers, which ensures that strong constraints are always satisfied in preference to weak ones.

Instances of `CISymbolicWeight` are used to represent these symbolic weights. These instances have an array of floating-point numbers, whose length is the number of nonrequired strengths (3 at the moment). Each element of the array represents the value at that strength, so [1.0, 0.0, 10.0] represents a weight of 1.0 strong, 0.0 medium, and 10.0 weak. (In Smalltalk, `CISymbolicWeight` is a variable-length subclass; we could have had an instance variable with an array of length 3 instead.) Symbolic weights understand various arithmetic messages (or operator overloading in C++), as follows:

`+ w.`

w is also a symbolic weight. Return the result of adding w to self (or this in C++).

`- w.`

w is also a symbolic weight. Return the result of subtracting w from self.

`*n.`

n is a number. Return the result of multiplying self by n.

`/n.`

n is a number. Return the result of dividing self by n.

`<= n, >= n, < n, > n, = n`

w is a symbolic weight. Return true if self is related to n as the operator normally queries.

`negative.`

Return true if this symbolic weight is negative (i.e., it does not consist of all zeros and the first nonzero number is negative).

Instances of `CIStructure` represent a strength in the constraint hierarchy. The instance variables are `name` (for printing purposes) and `symbolicWeight`, which is the unit symbolic weight for this strength. Thus, with the 3 strengths as above, `strong` is [1.0, 0.0, 0.0], `medium` is [0.0, 1.0, 0.0], and `weak` is [0.0, 0.0, 1.0].

The above arithmetic messages let the Smalltalk implementation of the solver use symbolic weights just like numbers in expressions. This interface is important because the objective row in the tableau has coefficients that are `CISymbolicWeights`, but are subject to the same manipulations as the other tableau rows whose expressions have coefficients that are just real numbers.

In both C++ and Java, an additional message `asDouble()` is understood by `CISymbolicWeights`. This converts the representation to a real number that approximates the total ordering suggested by the more general vector of real numbers. It is these real numbers that are used as the coefficients in the objective row of the tableau, instead of `CISymbolicWeights` (which conceptually the coefficients are). This kludge avoids the complexities that such genericity introduces to the static type systems of C++ and Java.

Also, since Java lacks operator overloading, the above operations are invoked using suggestive alphabetic method names such as `add`, `subtract`, `times`, and `lessThan`.

A.3 CISimplexSolver Implementation

Here are the instance variables of `CISimplexSolver` (some fields are inherited from `CITableau`, the base class of `CISimplexSolver`, which provides the basic sparse-matrix interface; see Section A.3.1).

`rows`. A dictionary with keys `CIAbstractVariable` and values `CILinearExpression`. This holds the tableau. Note that the keys can be either restricted or unrestricted variables, i.e., both C_U and C_S are actually merged into one tableau. This simplified the code considerably, since most operations are applied to both restricted and unrestricted rows.

`columns`. A dictionary with keys `CIAbstractVariable` and values `Set of CIAbstractVariable`. These are the column cross-indices. Each parametric variable p should be a key in this dictionary. The corresponding set should include exactly those basic variables whose linear expression includes p (p will of course have a nonzero coefficient). The keys can be either unrestricted or restricted variables.

`objective`. An instance of `CIObjectiveVariable` (named z) that is the key for the objective row in the tableau.

`infeasibleRows`. A set of basic variables that have infeasible rows. (This field is used when reoptimizing with the dual simplex method.)

`prevEditConstants`. An array of constants (floats) for the edit constraints on the previous iteration. The elements in this array must be in the same order

as `editPlusErrorVars` and `editMinusErrorVars`, and the argument to the public `resolve:` message.

`stayPlusErrorVars`, `stayMinusErrorVars`. An array of plus/minus error variables (instances of `CISlackVariable`) for the stay constraints. The corresponding negative/positive error variable must have the same index in `stayMinusErrorVars/stayPlusErrorVars`.

`editPlusErrorVars`, `editMinusErrorVars`. An array of plus/minus error variables (instances of `CISlackVariable`) for the edit constraints. The corresponding negative/positive error variable must have the same index in `editMinusErrorVars/editPlusErrorVars`.

`markerVars`. A dictionary whose keys are constraints and whose values are instances of a subclass of `CIAbstractVariable`. This dictionary is used to find the marker variable for a constraint when deleting that constraint. A secondary use is that iterating through the keys will give all of the original constraints (useful for the `reset` method).

`errorVars`. A dictionary whose keys are constraints and whose values are arrays of `CISlackVariable`. This dictionary gives the error variable (or variables) for a given nonrequired constraint. We need this if the constraint is deleted because the corresponding error variables must be deleted from the objective function.

A.3.1 CITableView (*Sparse Matrix*) Operations. The basic requirements for the tableau representation are that we should be able to perform the following operations efficiently:

- determine whether a variable is basic or parametric;
- find the corresponding expression for a basic variable;
- iterate through all the parametric variables with nonzero coefficients in a given row;
- find all the rows that contain a given parametric variable with a nonzero coefficient;
- add/remove a row;
- remove a parametric variable;
- substitute out a variable (i.e., replace all occurrences of a variable with an expression, updating the tableau as appropriate).

The representation of the tableau as a dictionary of rows, with column cross-indices, supports these operations. Keeping the cross indices up-to-date and consistent with the row dictionary is error-prone. Thus, the solver actually accesses the rows and columns only via the interface of `CITableView` below.

`addRow(CIAbstractVariable var, CILinearExpression expr)`. Add the constraint `var = expr` to the tableau. `var` will become a basic variable. Update the column cross indices.

`noteAddedVariable(CIAbstractVariable var, CIAbstractVariable subject)`. Variable `var` has been added to the linear expression for `subject`. Update the column cross indices.

`noteRemovedVariable(CIAbstractVariable var, CIAbstractVariable subject).`
 Variable `var` has been removed from the linear expression for `subject`. Update the column cross indices.

`removeColumn(CIAbstractVariable var)`. Remove the parametric variable `var` from the tableau. This operation involves removing the column cross index for `var` and removing `var` from every expression in `rows` in which it occurs.

`removeRow(CIAbstractVariable var)`. Remove the basic variable `var` from the tableau. Because `var` is basic, there should be a row `var = expr`. Remove this row, and also update the column cross indices.

`substituteOut(CIAbstractVariable var, CILinearExpression expr)`. Replace all occurrences of `var` with `expr` and update the column cross indices.

ACKNOWLEDGMENTS

Thanks to Kim Marriott for his work on the closely-related QOCA algorithm and comments on Cassowary. We received much useful feedback from early users of our toolkit, including Anthony Beurivé, Alexandre Duret-Lutz, Michael Kaufmann, Brian Grant, Pengling He, Tessa Lau, Sorin Lerner, John MacPhail, Larry Melia, Michael Noth, Emmanuel Pietriga, Stefan Saroiu, and Steve Wolfman.

REFERENCES

- BADROS, G., TIRTOWIDJOJO, J. J., MARRIOTT, K., MEYER, B., PORTNOY, W., AND BORNING, A. 2001. A constraint extension to scalable vector graphics. In *Proceedings of the Tenth World Wide Web Conference, WWW'10*.
- BADROS, G. J. AND BORNING, A. 2001. Cassowary constraint solving toolkit. Web page. <http://www.cs.washington.edu/research/constraints/cassowary>.
- BADROS, G. J., BORNING, A., MARRIOTT, K., AND STUCKEY, P. 1999. Constraint cascading style sheets for the web. In *Proceedings of the 1999 ACM Conference on User Interface Software and Technology*. ACM, New York, NY.
- BADROS, G. J., NICHOLS, J., AND BORNING, A. 2000. SCWM—the Scheme Constraints Window Manager. In *Proceedings of the AAAI Spring Symposium on Smart Graphics*.
- BARAFF, D. 1994. Fast contact force computation for nonpenetrating rigid bodies. In *SIGGRAPH '94 Conference Proceedings*. ACM, New York, NY, 23–32.
- BORNING, A., ANDERSON, R., AND FREEMAN-BENSON, B. 1996. Indigo: A local propagation algorithm for inequality constraints. In *Proceedings of the 1996 ACM Symposium on User Interface Software and Technology*. ACM, New York, NY, 129–136.
- BORNING, A. AND BADROS, G. 2000. On finding graphically plausible solutions to constraint hierarchies: The split stay problem. In *Workshop on Soft Constraints: Theory and Practice*. Sixth International Conference on Principles and Practice of Constraint Programming (Singapore). Available from <http://www.cs.washington.edu/research/constraints>.
- BORNING, A. AND FREEMAN-BENSON, B. 1998. Ultraviolet: A constraint satisfaction algorithm for interactive graphics. *Constraints* 3, 1 (Apr.), 9–32.
- BORNING, A., FREEMAN-BENSON, B., AND WILSON, M. 1992. Constraint hierarchies. *Lisp and Symbolic Computation* 5, 3 (Sept.), 223–270.
- BORNING, A., LIN, R., AND MARRIOTT, K. 1997. Constraints for the web. In *Proceedings of ACM MULTIMEDIA '97*. ACM, New York, NY.
- BORNING, A., LIN, R., AND MARRIOTT, K. 2000. Constraint-based document layout for the web. *Multimedia Syst.* 8, 3, 177–189.
- BORNING, A., MARRIOTT, K., STUCKEY, P., AND XIAO, Y. 1997. Solving linear arithmetic constraints for user interface applications. In *Proceedings of the 1997 ACM Symposium on User Interface Software and Technology*. ACM, New York, NY.

- HOSOBE, H., MATSUOKA, S., AND YONEZAWA, A. 1996. Generalized local propagation: A framework for solving constraint hierarchies. In *Proceedings of the Second International Conference on Principles and Practice of Constraint Programming*. LINCS 1118, Springer-Verlag, Heidelberg, Germany.
- HUDSON, S. AND SMITH, I. 1996. SubArctic UI toolkit user's manual. Tech. Rep., College of Computing, Georgia Institute of Technology.
- HUYNH, T. AND MARRIOTT, K. 1995. Incremental constraint deletion in systems of linear constraints. *Inf. Process. Lett.* 55, 111–115.
- JAFFAR, J., MICHA耶LOV, S., STUCKEY, P., AND YAP, R. 1992. The CLP(\mathcal{R}) language and system. *ACM Trans. Program. Lang. Syst.* 14, 3 (July), 339–395.
- MARRIOTT, K. AND STUCKEY, P. 1998. *Programming with Constraints: An Introduction*. MIT Press, Cambridge, MA.
- MYERS, B. A. 1996. The Amulet user interface development environment. In *CHI'96 Conference Companion: Human Factors in Computing Systems*. ACM, New York, NY.
- PRESS, W. H., FLANNERY, B. P., TEUKOLSKY, S. A., AND VETTERLING, W. T. 1989. *Numerical Recipes: The Art of Scientific Computing*, 2nd ed. Cambridge University Press, Cambridge, MA.
- SANNELLA, M. 1994. SkyBlue: A multi-way local propagation constraint solver for user interface construction. In *Proceedings of the 1994 ACM Symposium on User Interface Software and Technology*. ACM, New York, NY, 137–146.
- SANNELLA, M., MALONEY, J., FREEMAN-BENSON, B., AND BORNING, A. 1993. Multi-way versus one-way constraints in user interfaces: Experience with the DeltaBlue algorithm. *Softw. Pract. Exper.* 23, 5 (May), 529–566.
- SUTHERLAND, I. 1963. Sketchpad: A man-machine graphical communication system. In *Proceedings of the Spring Joint Computer Conference*. IFIPS, 329–346.
- VAN HENTENRYCK, P. AND GRAF, T. 1990. Standard forms for rational linear arithmetic in constraint logic programming. Internal Rep. ir-lp-2217, European Computer Research Centre.
- VANDER ZANDEN, B. 1996. An incremental algorithm for satisfying hierarchies of multi-way dataflow constraints. *ACM Trans. Program. Lang. Syst.* 18, 1 (Jan.), 30–72.
- WOLFMAN, S. AND WELD, D. 1999. The LPSAT engine and its application to resource planning. In *Proceedings of IJCAI'99*.

Received January 2000; July 2001; accepted November 2001